

Chapter 1: roadmap

1.1 What is the Internet?

1.2 Network edge

- ❖ end systems, access networks, links

1.3 Network core

- ❖ circuit switching vs packet switching,
- ❖ network structure

1.4 Performance:

- ❖ delay, loss and throughput

1.5 Protocol layers, service models

1.6 Networks under attack: security

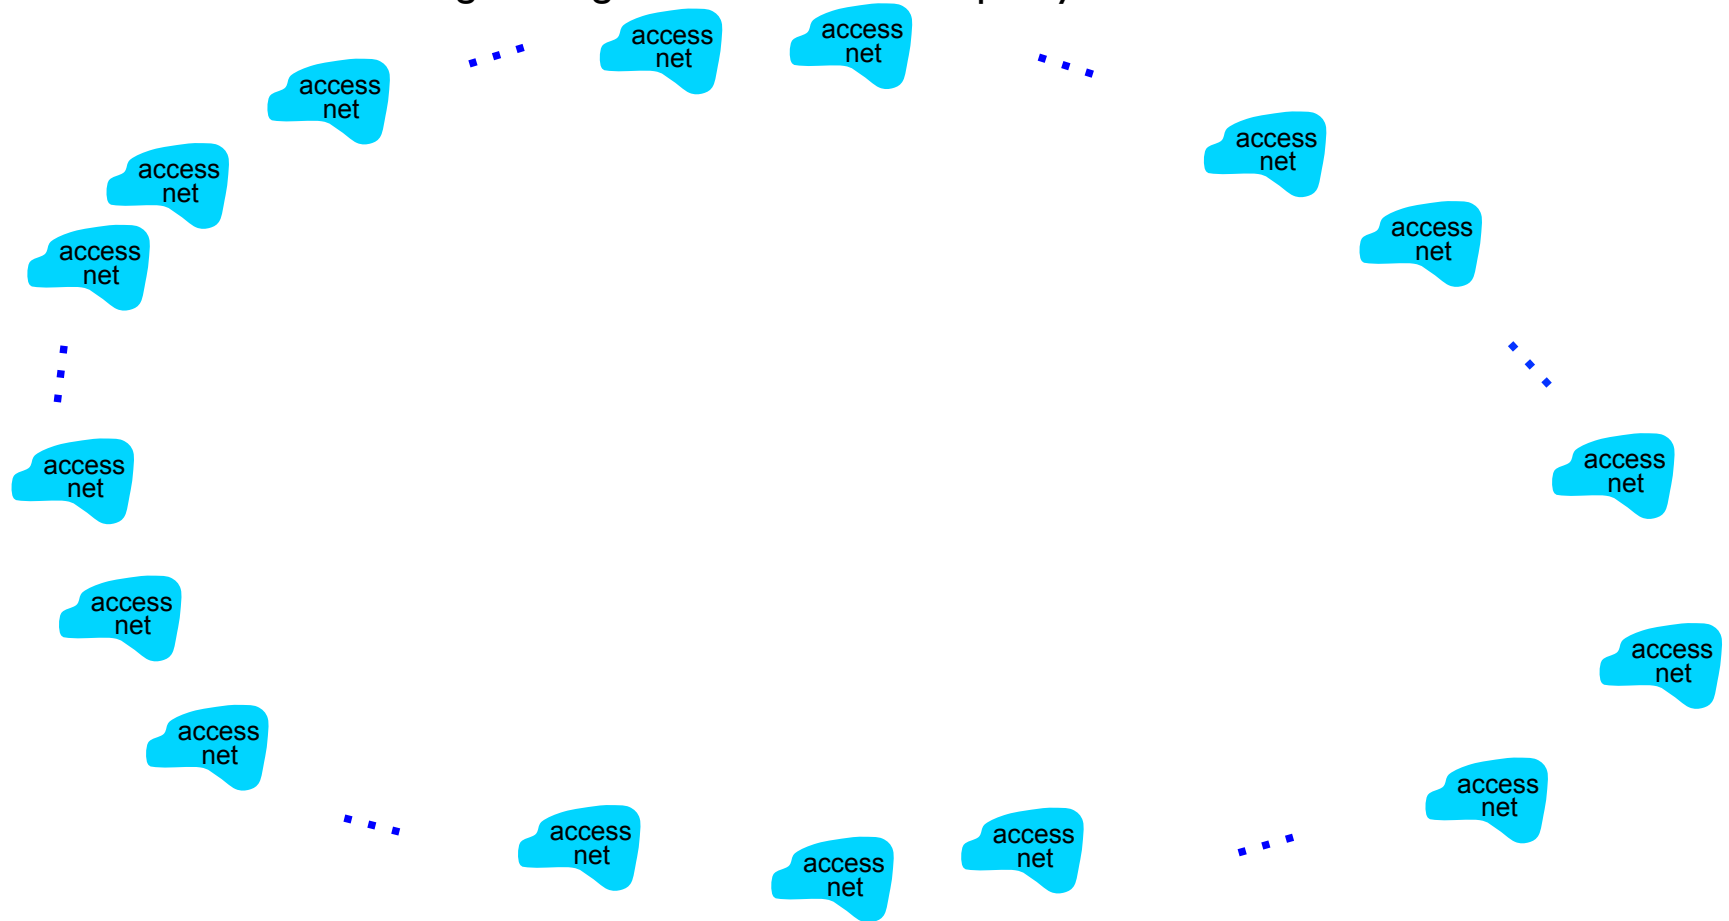
1.7 History

Internet Structure: network of networks

- ❖ End systems connect to Internet via **access ISPs** (Internet Service Providers)
 - Residential, company and university ISPs
- ❖ Access ISPs in turn must be interconnected.
 - ❖ So that any two hosts can send packets to each other
- ❖ Resulting network of networks is very complex
 - ❖ Evolution was driven by **economics** and **national policies**
- ❖ Let's take a stepwise approach to describe current Internet structure

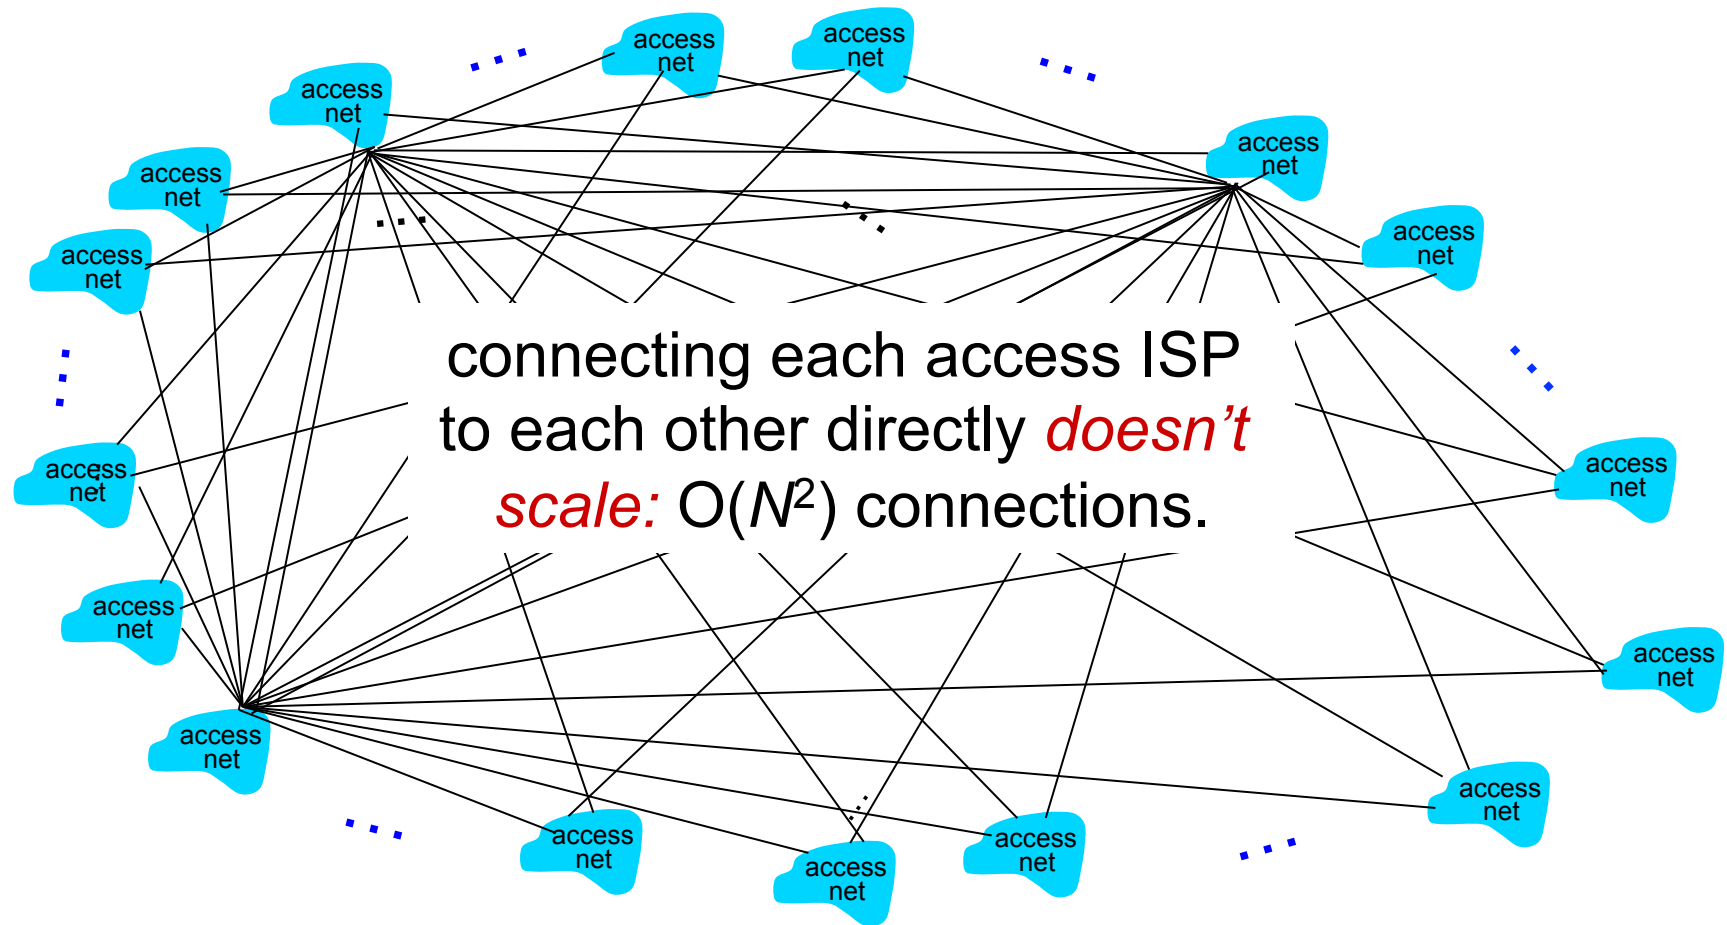
Internet structure: network of networks

- ❖ End systems connect to Internet via **access ISPs**:
 - ❖ residential, company and university ISPs
- ❖ **Question:** given *millions* of access ISPs, how to connect them together?
 - ❖ Considerations: engineering but also economics+policy




Internet structure: network of networks

Option: connect each access ISP to every other access ISP?



Internet structure: network of networks


*Option: connect each access ISP to a global transit ISP? **Customer** and **provider** ISPs have economic agreement.*



Internet structure: network of networks

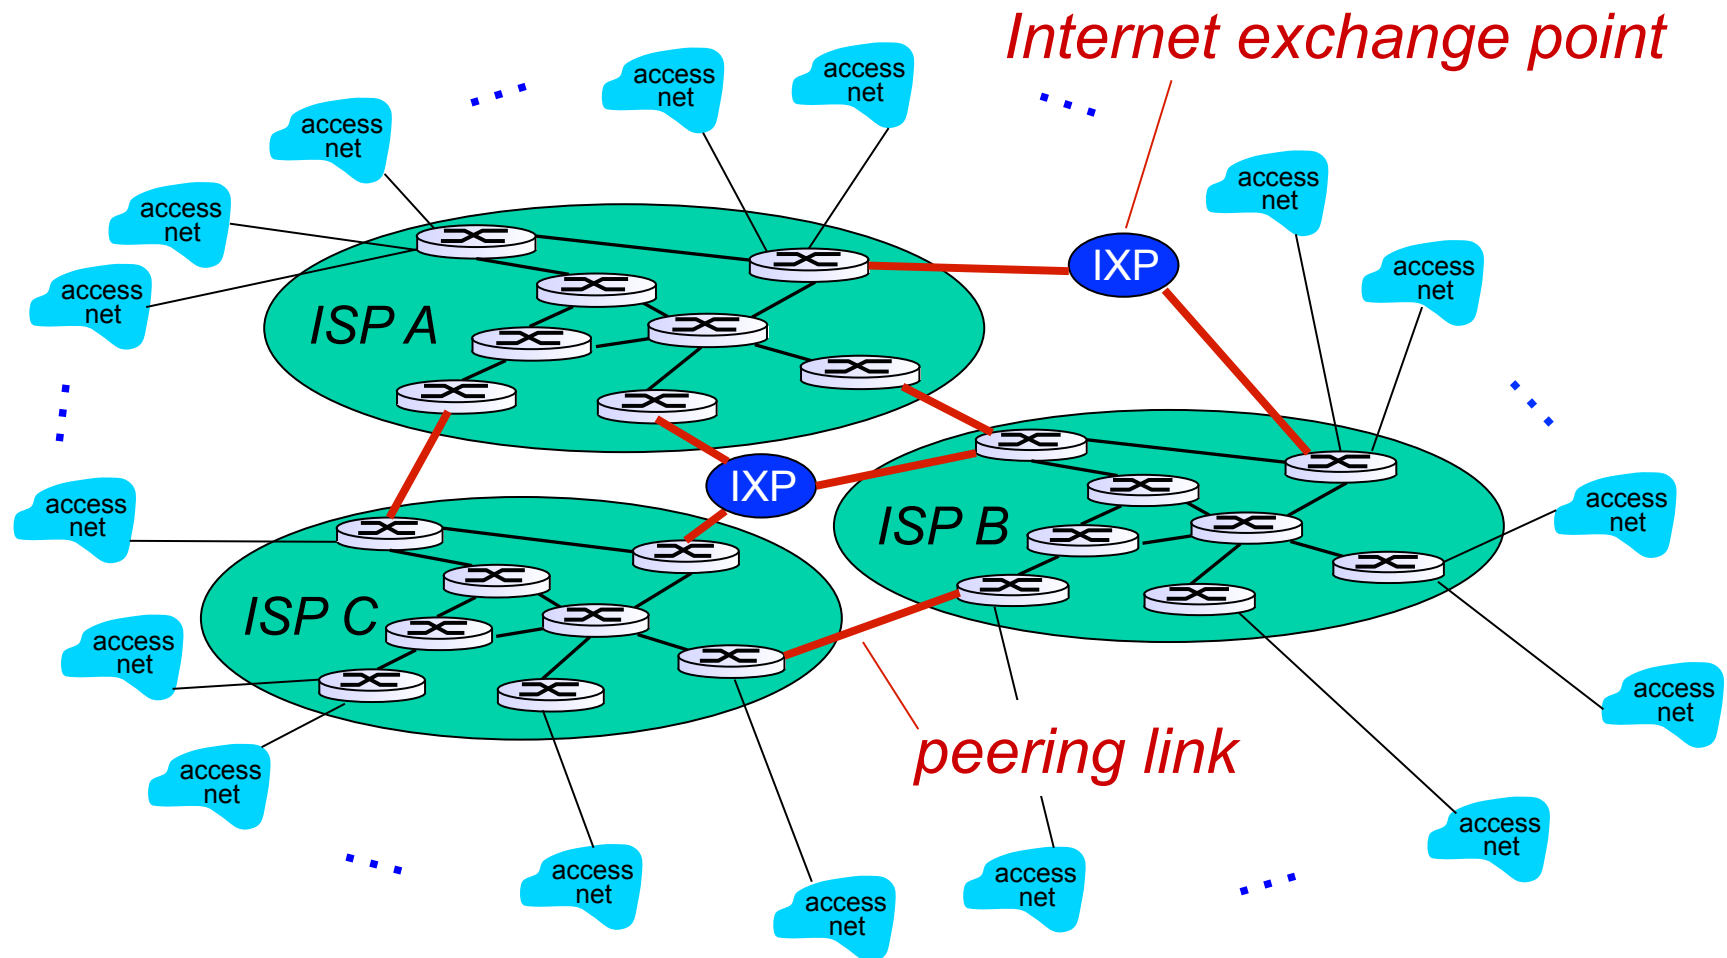
But if one global ISP is viable business, there will be competitors

....



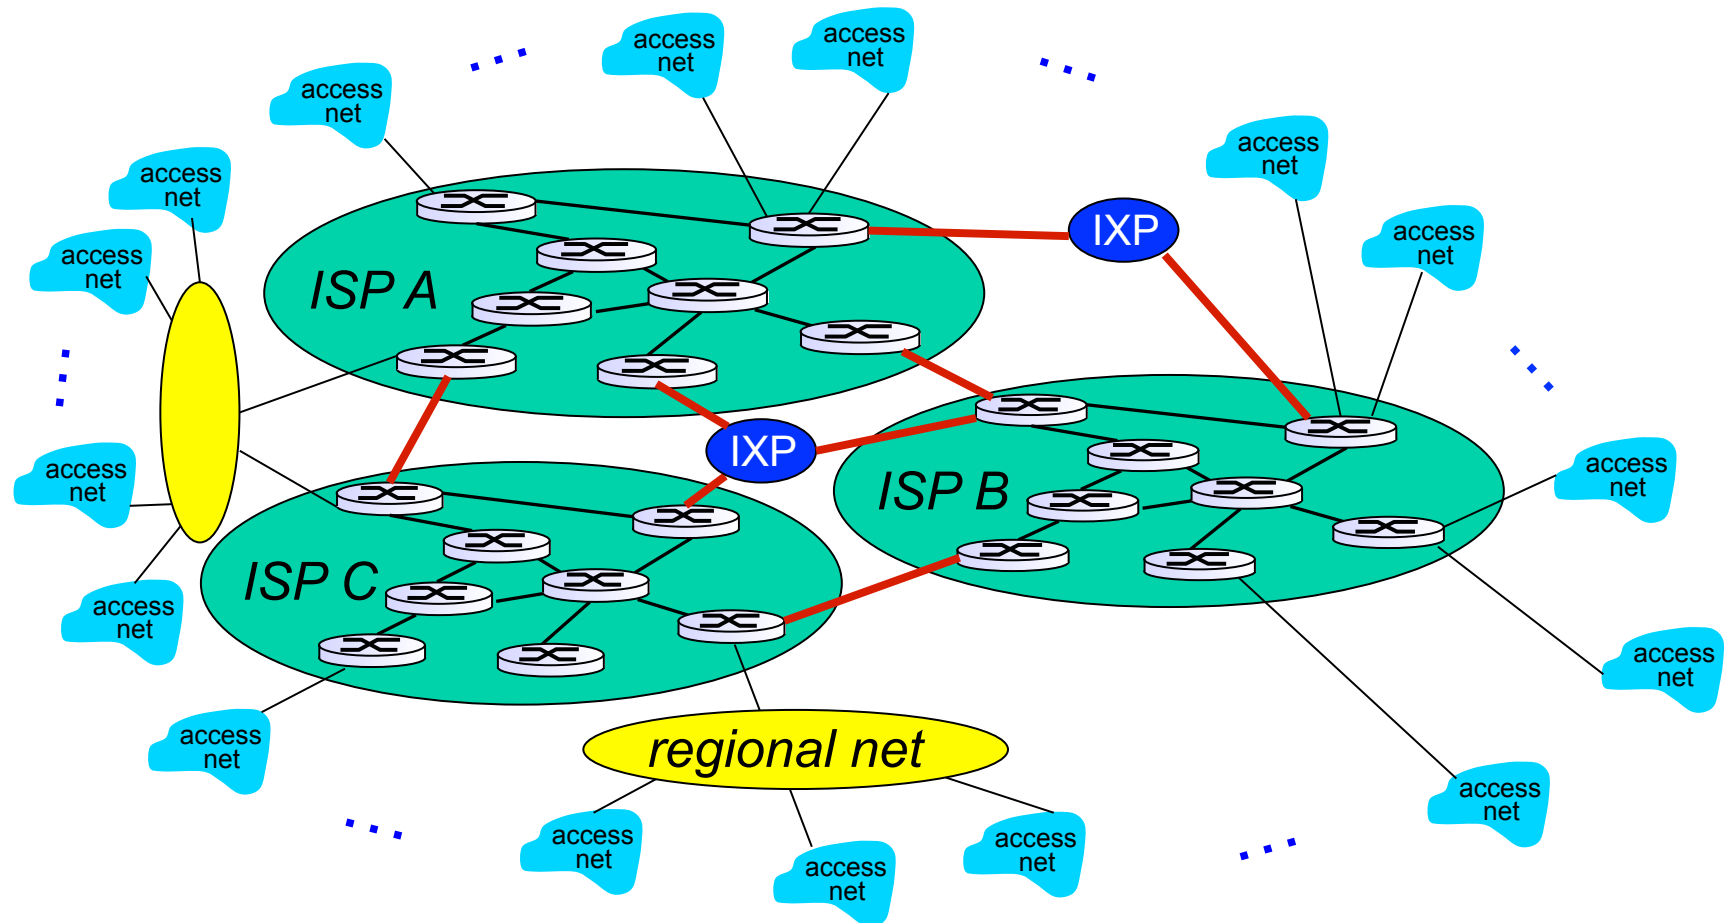
Internet structure: network of networks

But if one global ISP is viable business, there will be competitors
.... which must be interconnected



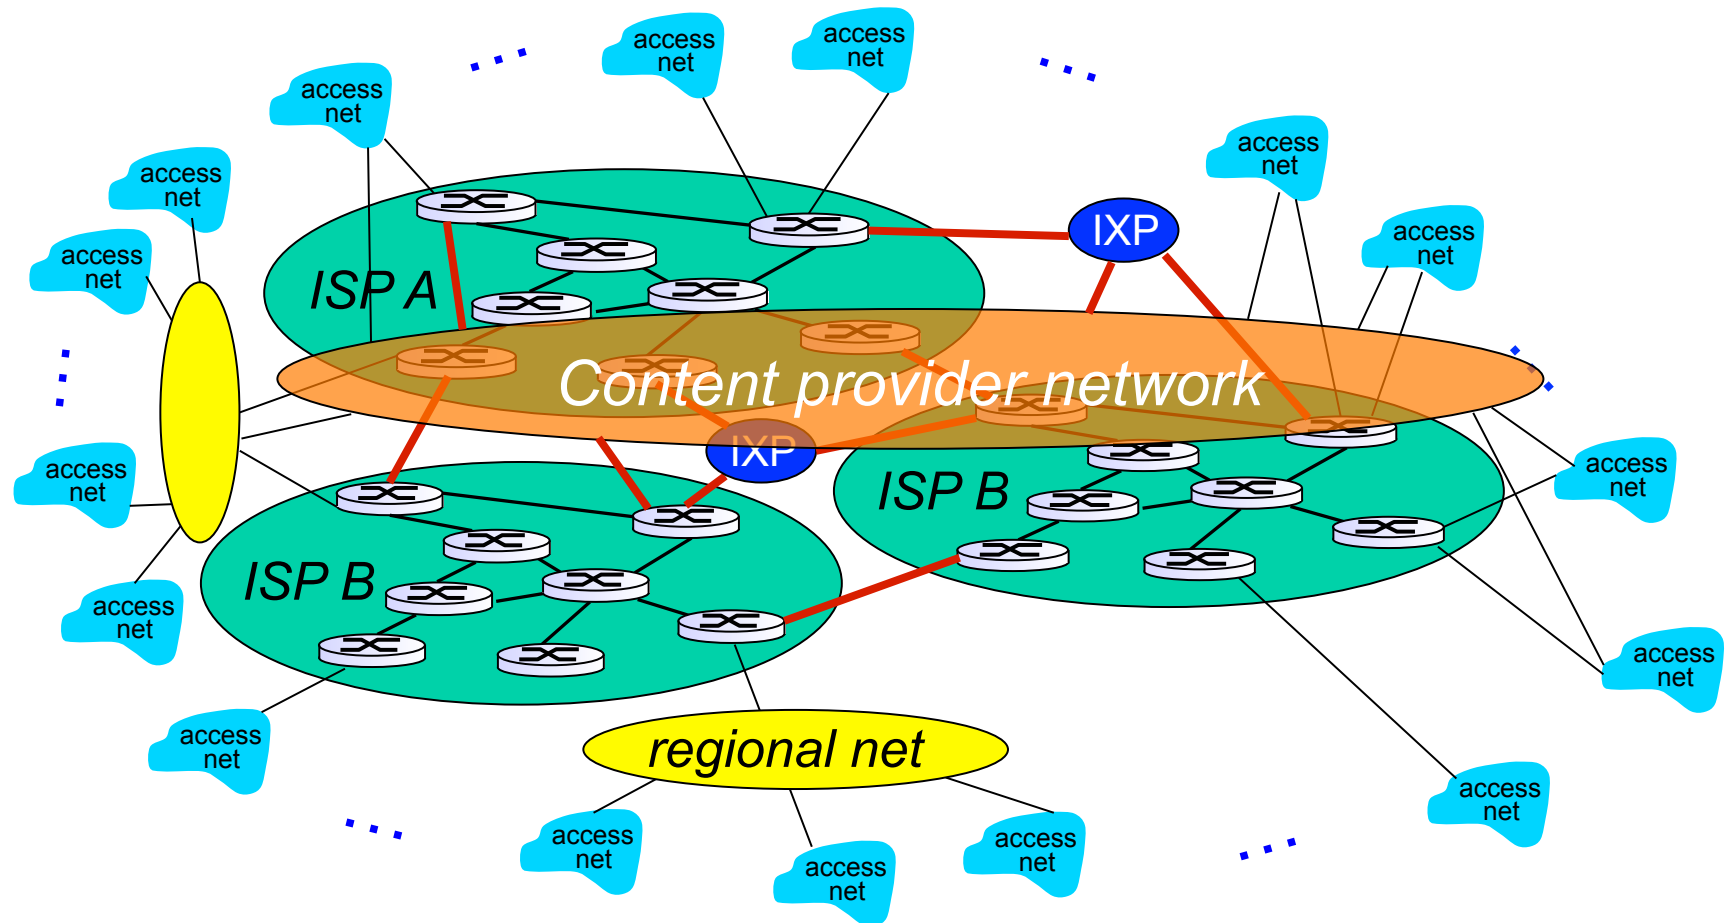
Internet structure: network of networks

... and regional networks may arise to connect access nets to ISPs

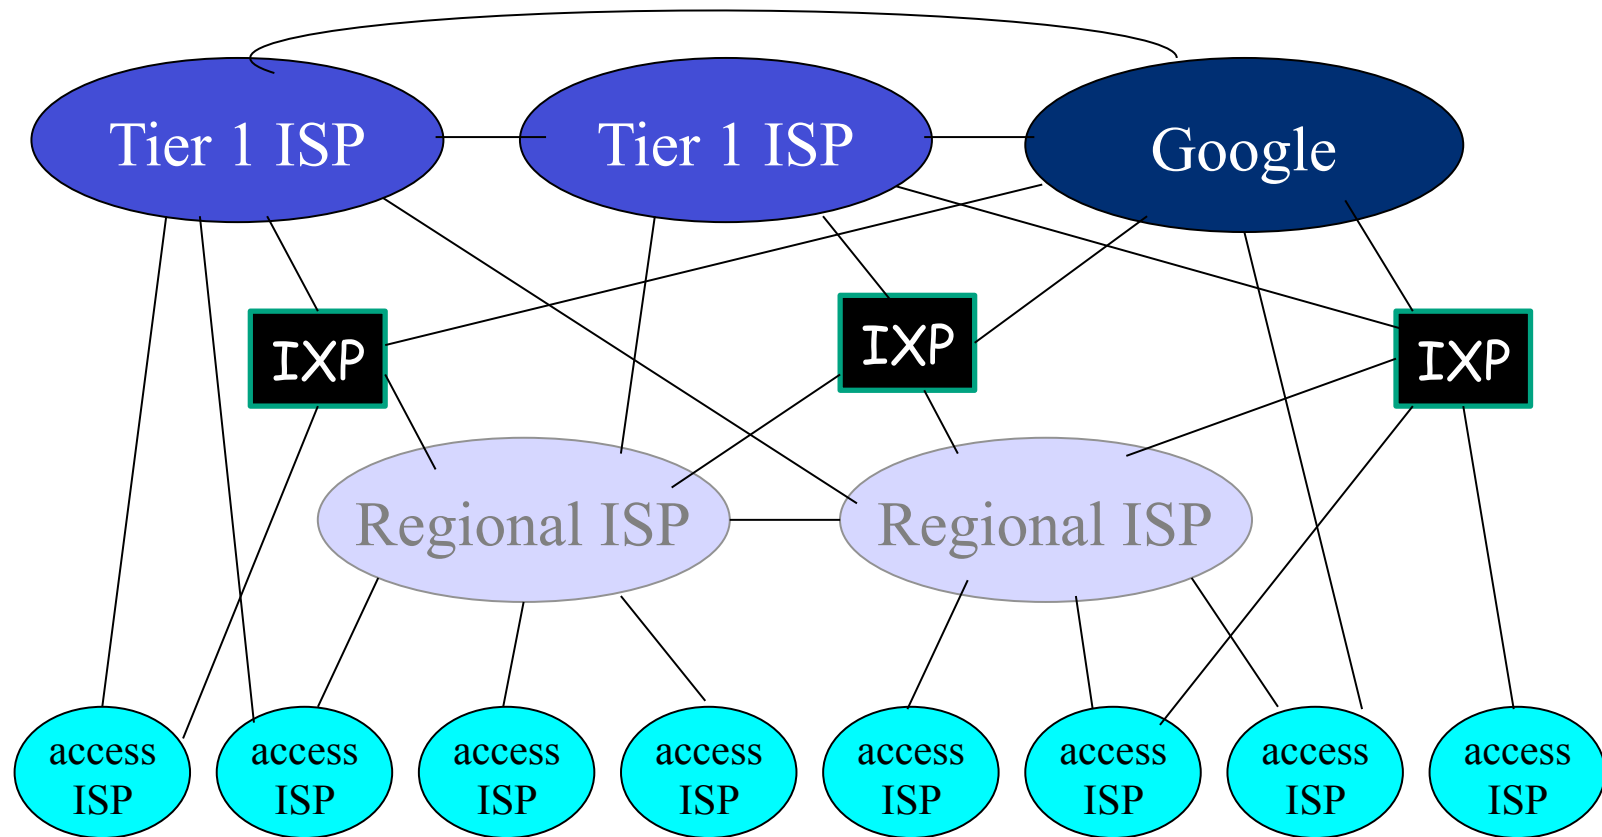


Internet structure: network of networks

... and content provider networks (e.g., Google, Microsoft, Akamai) may run their own network, to bring services, content close to end users

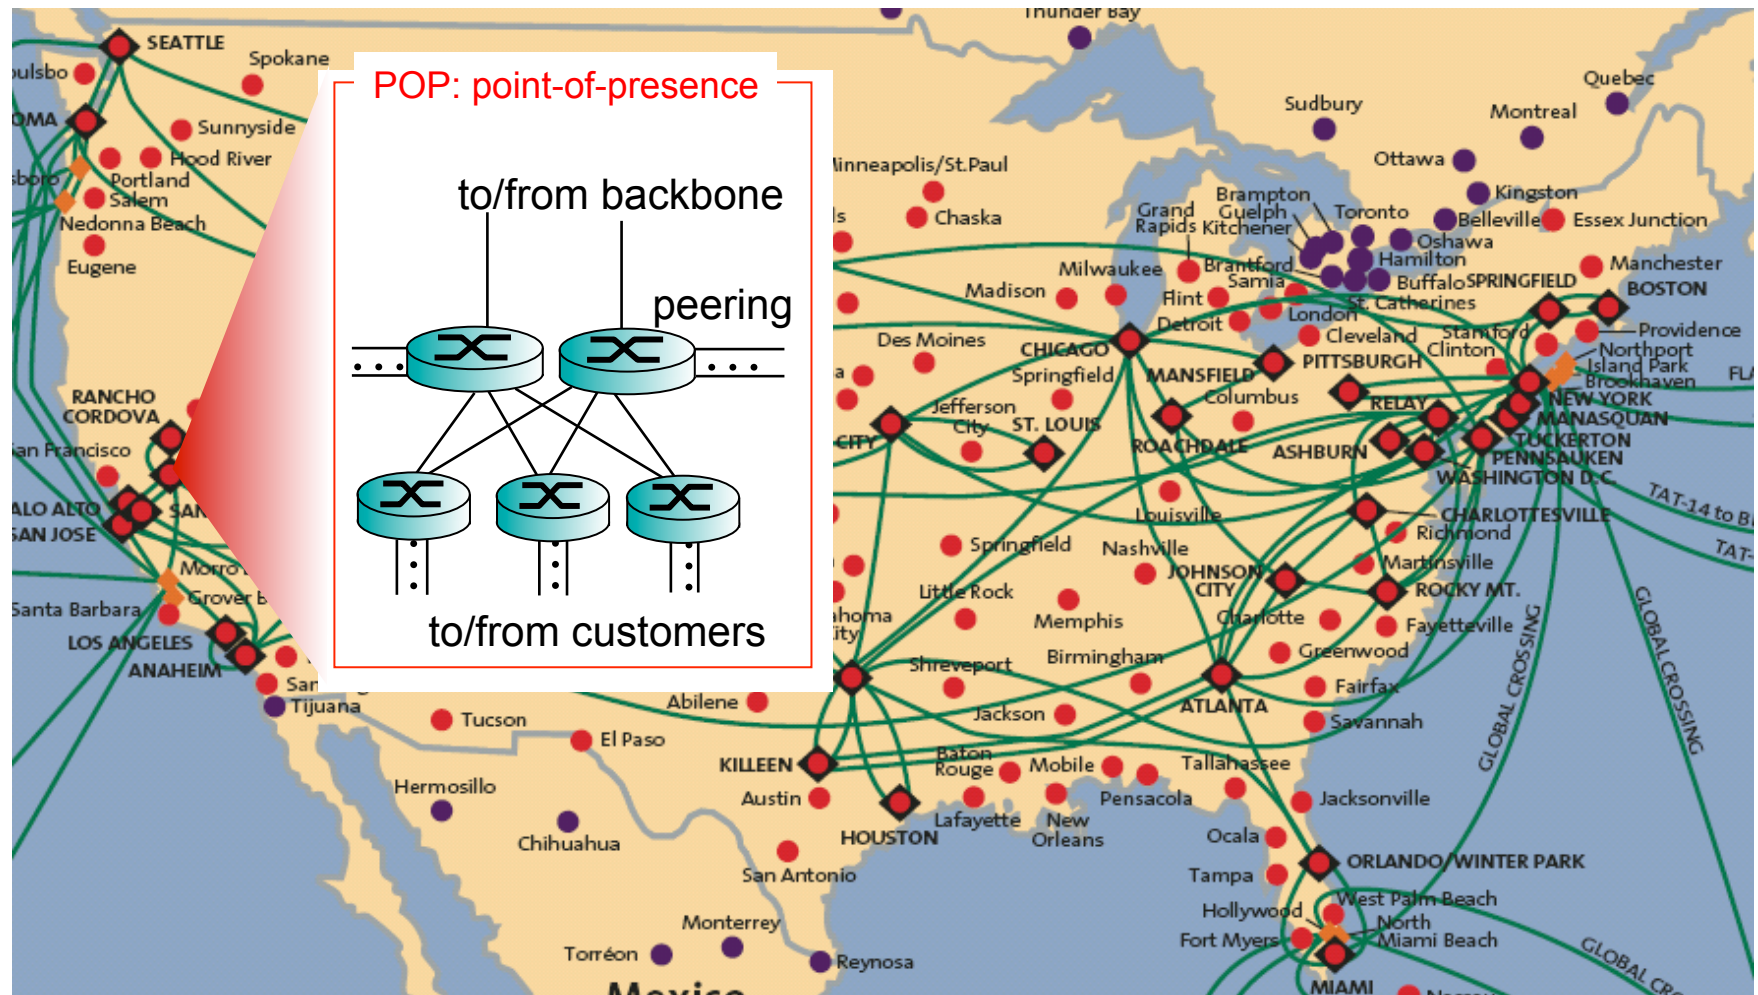


Internet structure: network of networks




- ❑ at center: small # of well-connected large networks
 - ❖ “tier-1” commercial ISPs (e.g., Level 3, Sprint, AT&T, NTT), national & international coverage
 - ❖ content provider network (e.g., Google): private network that connects its data centers to Internet, often bypassing tier-1, regional ISPs

Tier-I ISP: e.g., Sprint



“Real” Internet delays and routes


- ❑ what do “real” Internet delay & loss look like?
- ❑ `traceroute` program: provides delay measurement from source to router along end-end Internet path towards destination. For all i :
 - ❖ sends three packets that will reach router i on path towards destination
 - ❖ router i will return packets to sender
 - ❖ sender times interval between transmission and reply.



“Real” Internet delays and routes


traceroute: from gaia.cs.umass.edu to www.eurecom.fr

Three delay measurements from
gaia.cs.umass.edu to cs-gw.cs.umass.edu




```
1 cs-gw (128.119.240.254) 1 ms 1 ms 2 ms
2 border1-rt-fa5-1-0.gw.umass.edu (128.119.3.145) 1 ms 1 ms 2 ms
3 cht-vbns.gw.umass.edu (128.119.3.130) 6 ms 5 ms 5 ms
4 jnl-atl-0-0-19.wor.vbns.net (204.147.132.129) 16 ms 11 ms 13 ms
5 jnl-so7-0-0-0.wae.vbns.net (204.147.136.136) 21 ms 18 ms 18 ms
6 abilene-vbns.abilene.ucaid.edu (198.32.11.9) 22 ms 18 ms 22 ms
7 nycm-wash.abilene.ucaid.edu (198.32.8.46) 22 ms 22 ms 22 ms
8 62.40.103.253 (62.40.103.253) 104 ms 109 ms 106 ms
9 de2-1.del.de.geant.net (62.40.96.129) 109 ms 102 ms 104 ms
10 de.fr1.fr.geant.net (62.40.96.50) 113 ms 121 ms 114 ms
11 renater-gw.fr1.fr.geant.net (62.40.103.54) 112 ms 114 ms 112 ms
12 nio-n2.cssi.renater.fr (193.51.206.13) 111 ms 114 ms 116 ms
13 nice.cssi.renater.fr (195.220.98.102) 123 ms 125 ms 124 ms
14 r3t2-nice.cssi.renater.fr (195.220.98.110) 126 ms 126 ms 124 ms
15 eurecom-valbonne.r3t2.ft.net (193.48.50.54) 135 ms 128 ms 133 ms
16 194.214.211.25 (194.214.211.25) 126 ms 128 ms 126 ms
17 ***
18 ***
19 fantasia.eurecom.fr (193.55.113.142) 132 ms 128 ms 136 ms
```

trans-oceanic link



* means no response (probe lost, router not replying)



More traceroute examples

Athina-Markopoulous-MacBook-Air-2:web athina\$ **traceroute gaia.cs.umass.edu**

traceroute to gaia.cs.umass.edu (128.119.245.12), 64 hops max, 52 byte packets

```
1 192.168.0.1 (192.168.0.1) 1.546 ms 0.247 ms 0.618 ms
2 10.75.152.1 (10.75.152.1) 6.692 ms 7.636 ms 7.960 ms
3 ip68-4-11-142.oc.oc.cox.net (68.4.11.142) 14.459 ms 8.553 ms 7.195 ms
4 ip68-4-11-12.oc.oc.cox.net (68.4.11.12) 26.645 ms 8.021 ms 7.628 ms
5 langbprj01-ae2.0.rd.la.cox.net (68.1.0.136) 8.452 ms 9.666 ms 9.866 ms
6 lsan0.tr-cps.internet2.edu (206.223.123.199) 9.921 ms 11.280 ms 10.319 ms
7 xe-0-2-0.0.sttl0.tr-cps.internet2.edu (64.57.20.223) 34.258 ms 34.985 ms 34.152 ms
8 xe-1-1-0.0.chic0.tr-cps.internet2.edu (64.57.20.169) 87.729 ms 88.794 ms 88.357 ms
9 198.71.47.62 (198.71.47.62) 112.777 ms 111.206 ms 111.823 ms
10 192.5.89.21 (192.5.89.21) 98.682 ms 98.637 ms 97.918 ms
11 nox300gw1-peer--207-210-142-242.nox.org (207.210.142.242) 98.917 ms 100.046 ms 99.847 ms
12 core1-rt-xe-0-0-0.gw.umass.edu (192.80.83.101) 100.235 ms 99.676 ms 98.883 ms
13 lgrc-rt-106-8-po-10.gw.umass.edu (128.119.0.233) 100.358 ms 100.123 ms 99.548 ms
14 128.119.3.32 (128.119.3.32) 101.201 ms 100.145 ms 99.514 ms
15 nscs1bbs1.cs.umass.edu (128.119.240.253) 101.423 ms 100.951 ms 103.533 ms
16 * * *
17 * * *
```

More traceroute examples

Athina-Markopoulous-MacBook-Air-2:web athina\$ **traceroute nestor.calit2.uci.edu**

traceroute to nestor.calit2.uci.edu (128.195.185.109), 64 hops max, 52 byte packets

```
1 302-wism-vl970.ucinet.uci.edu (169.234.0.1) 59.342 ms 27.768 ms 90.426 ms
2 cs1-core--302-wism.ucinet.uci.edu (128.195.249.129) 1.325 ms 1.361 ms 1.254 ms
3 325--cs1-core.ucinet.uci.edu (128.195.249.190) 1.051 ms 1.113 ms 1.035 ms
4 nestor.calit2.uci.edu (128.195.185.109) 1.019 ms 1.053 ms 1.030 ms
```

Athina-Markopoulous-MacBook-Air-2:web athina\$ **traceroute www.google.com**

traceroute: Warning: www.google.com has multiple addresses; using 74.125.224.178

traceroute to www.google.com (74.125.224.178), 64 hops max, 52 byte packets

```
1 302-wism-vl970.ucinet.uci.edu (169.234.0.1) 2.994 ms 0.930 ms 1.018 ms
2 cs1-core--302-wism.ucinet.uci.edu (128.195.249.129) 1.036 ms 1.003 ms 0.966 ms
3 kazad-dum-vl123.ucinet.uci.edu (128.200.2.222) 1.286 ms 1.088 ms 1.234 ms
4 dc-tus-agg2--uci-ge-1.cenic.net (137.164.24.49) 2.816 ms 2.306 ms 1.700 ms
5 riv-core1--tus-agg2-10ge-2.cenic.net (137.164.47.79) 9.473 ms 6.731 ms 13.066 ms
6 dc-lax-core1--riv-core1-10ge-2.cenic.net (137.164.46.57) 12.061 ms 9.388 ms 11.659 ms
7 72.14.223.85 (72.14.223.85) 7.697 ms 7.713 ms 7.751 ms
8 64.233.174.238 (64.233.174.238) 7.777 ms 9.668 ms 8.567 ms
9 72.14.236.11 (72.14.236.11) 8.098 ms 8.008 ms 8.128 ms
10 lax02s01-in-f18.1e100.net (74.125.224.178) 8.130 ms 7.953 ms 8.101 ms
```

Athina-Markopoulous-MacBook-Air-2:web athina\$

Chapter 1: roadmap

1.1 what is the Internet?

1.2 network edge

- end systems, access networks, links

1.3 network core

- packet switching, circuit switching, network structure

1.4 delay, loss, throughput in networks

1.5 protocol layers, service models

1.6 networks under attack: security

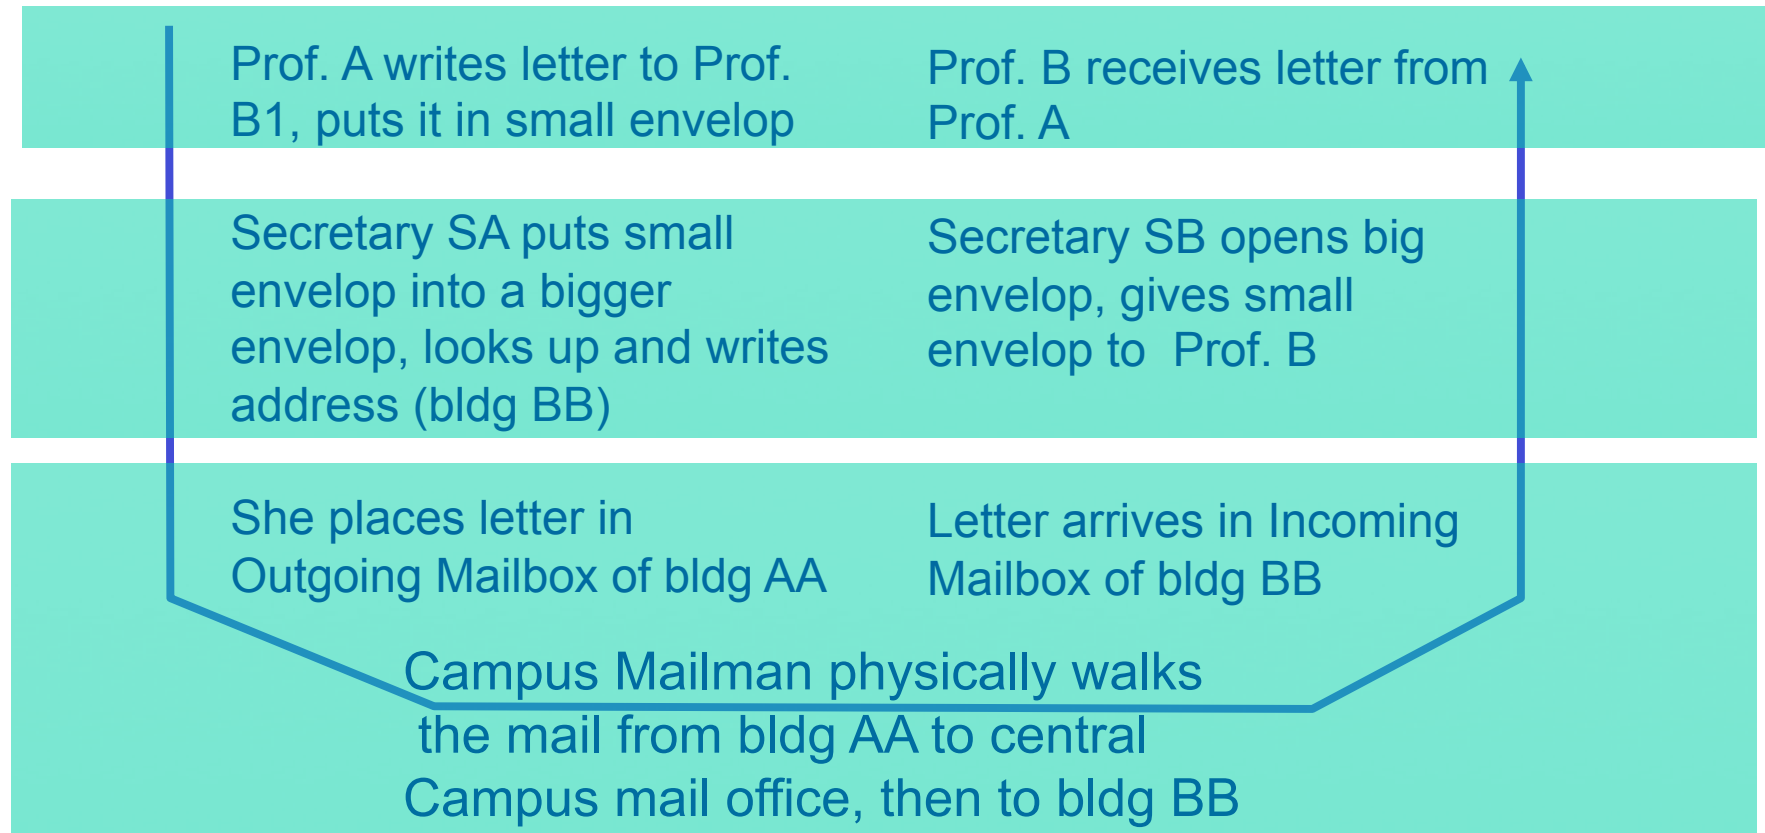
1.7 history

Dealing with Scale and Complexity

- ❑ The Internet is highly complex, in terms of ...
 - ❖ Its sheer **size** (number of components)
 - ❖ The number of **tasks** it needs to manage: routing, congestion control, packet reordering, connection establishment ...
 - ❖ Diversity of Components, Topology, Functionality,
- ❑ Tasks are structured through **modularization**
 - ❖ Divide and Conquer
 - ❖ Break down into **smaller pieces**
 - ❖ Create different functional layers



Layering of Campus Mail



layers: each layer implements a service

- ❖ via its own internal-layer actions
- ❖ relying on services provided by layer below

Why layering?

dealing with complex systems:

- ❑ explicit structure allows identification, relationship of complex system's pieces
 - ❖ layered *reference model* for discussion
- ❑ modularization eases maintenance, updating of system
 - ❖ change of implementation of layer's service transparent to rest of system
 - ❖ e.g., change in secretary, secretary's routine, mailman's routine, doesn't affect rest of system
- ❑ layering considered harmful?
 - ❖ duplication of functionality?
 - ❖ cross-layer optimization?

Internet protocol stack

❑ *Application layer:*

- ❖ What we interact with:

❑ *Transport layer:*

- ❖ process-process data transfer
- ❖ End-to-end management: TCP, UDP

❑ *Network layer:*


- ❖ routing of datagrams from source to destination
- ❖ IP, routing protocols

❑ *Link Layer:*

- ❖ data transfer between neighboring network elements
- ❖ Ethernet, 802.111 (WiFi), PPP


❑ *physical:*

- ❖ network medium (fiber, wireless,...)
- ❖ bits on the wire



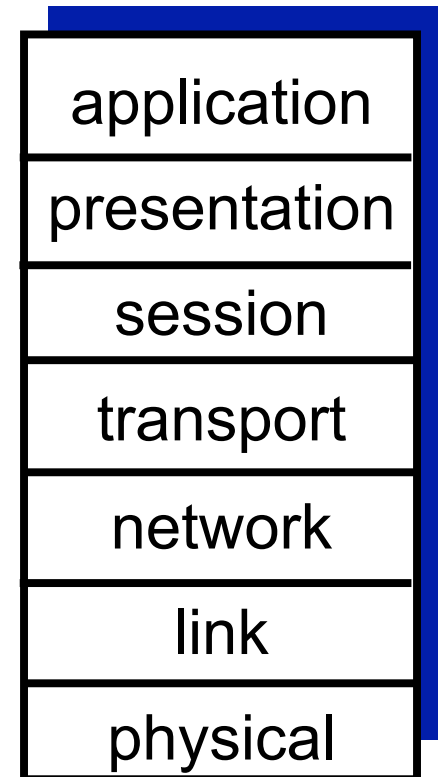
Transport & Network

- ❑ “Thin waist”
 - ❖ Other layers: evolved dramatically
 - ❖ TCP/IP: Stayed mostly the same
- ❑ Transport layer
 - ▶ End-to-end
 - ▶ Connection-oriented in TCP
- ❑ Network layer
 - ▶ Hop-by-hop
 - ▶ Connectionless in IP



ISO/OSI reference model

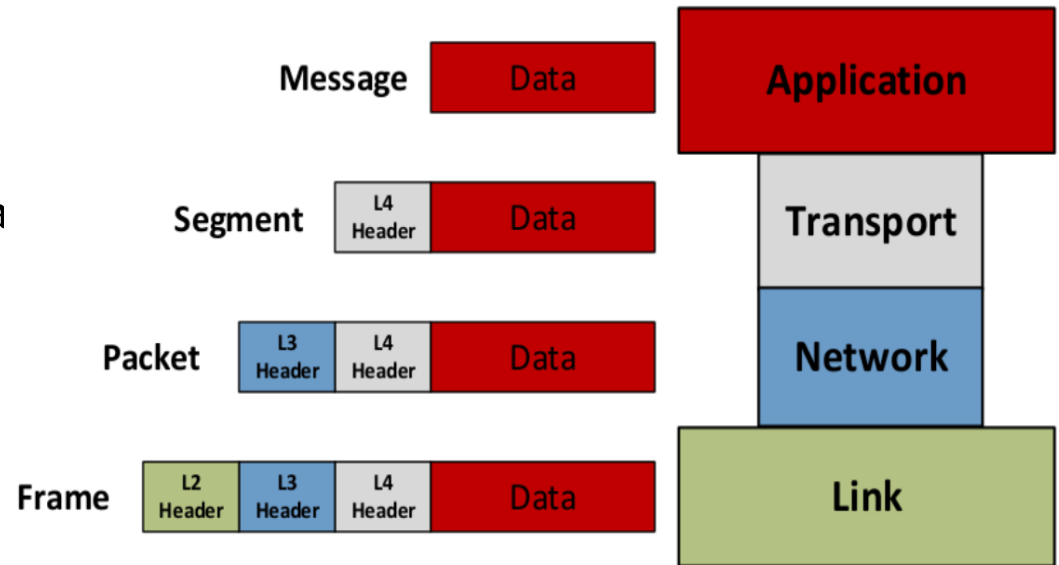
- ❑ *presentation*: allow applications to interpret meaning of data, e.g., encryption, compression, machine-specific conventions
- ❑ *session*: synchronization, checkpointing, recovery of data exchange
- ❑ Internet stack “missing” these layers!
 - ❖ these services, *if needed*, must be implemented in application



Headers

- ❑ **Encapsulation**: augment data with headers

- ❖ **Payload**: actual content
- ❖ **Header**: identification and control information



- ❑ This creates **overhead**. Why bother?
 - ▶ Allows us to distinguish messages
 - ▶ e.g.: Layer 3 header contains **src & dst IP address** of next hop
 - ▶ Essential for **packet switching**

“Understanding” Layers

- Different network elements process up to different layers

- End-hosts


- ❖ Server, computer
- ❖ Process **all 5**

- Routers

- ❖ **Up to 3**
- ❖ Have IP addresses

- Switches

- ❖ **Up to 2**
- ❖ Do not have or process IP



Encapsulation

