


# Chapter 3

## Transport Layer



*Computer  
Networking: A Top  
Down Approach*  
6<sup>th</sup> edition  
Jim Kurose, Keith Ross  
Addison-Wesley  
March 2012

© Based on materials developed by J.F.Kurose and K.W.Ross.  
All right reserved.

# Chapter 3: Transport Layer

## our goals:

- ❖ understand principles behind transport layer services:
  - multiplexing, demultiplexing
  - reliable data transfer
  - flow control
  - congestion control
- ❖ learn about Internet transport layer protocols:
  - UDP: connectionless transport
  - TCP: connection-oriented reliable transport
  - TCP congestion control

# Chapter 3 outline

## 3.1 transport-layer services

## 3.2 multiplexing and demultiplexing

## 3.3 connectionless transport: UDP

## 3.4 principles of reliable data transfer

## 3.5 connection-oriented transport: TCP

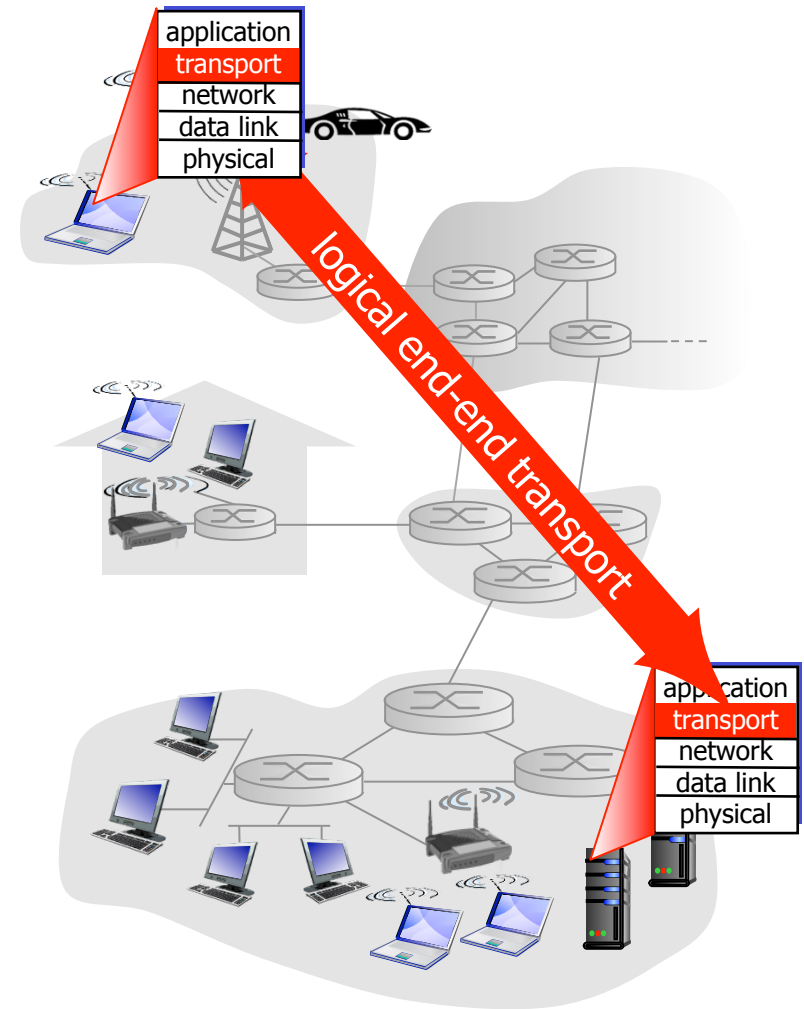
- segment structure
- reliable data transfer
- flow control
- connection management

## 3.6 principles of congestion control

## 3.7 TCP congestion control

# Transport services and protocols

- ❖ provide *logical communication* between app processes running on different hosts
- ❖ transport protocols run in end systems
  - send side: breaks app messages into *segments*, passes to network layer
  - rcv side: reassembles segments into messages, passes to app layer
- ❖ more than one transport protocol available to apps
  - Internet: TCP and UDP



# Transport vs. network layer

- ❖ *application layer*
- ❖ *transport layer*: logical communication between processes
  - relies on + enhances, network layer services
- ❖ *network layer*: logical communication between hosts

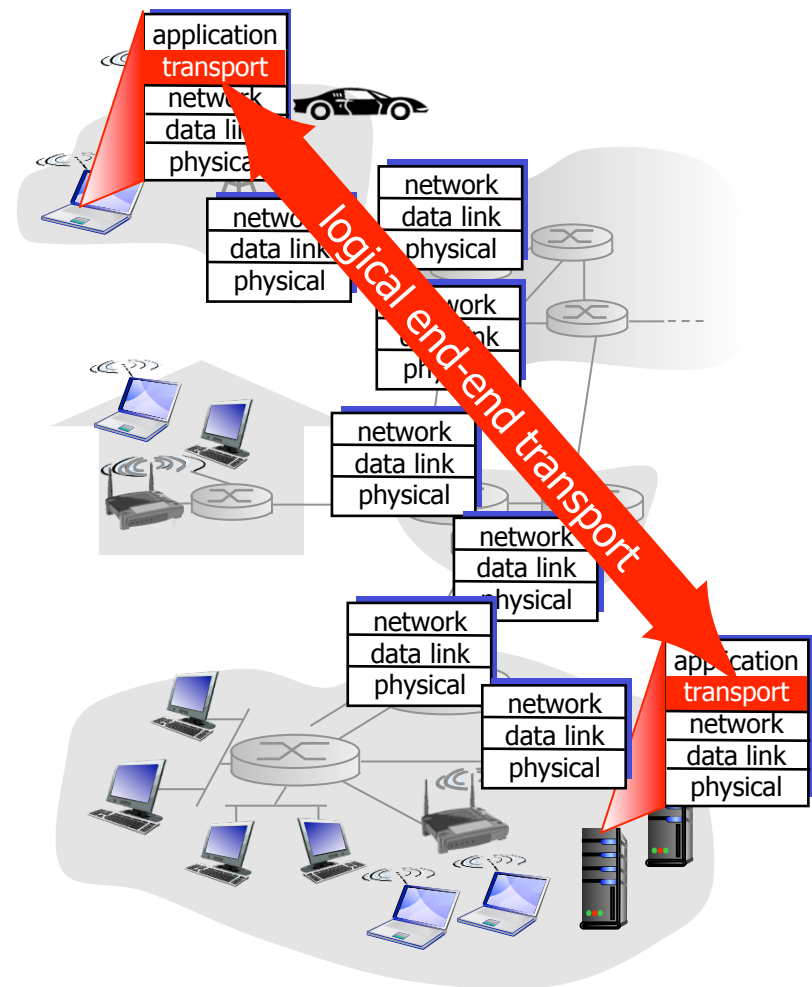
## *household analogy:*

12 kids in Ann's house sending letters to 12 kids in Bill's house:

- ❖ hosts = houses
- ❖ processes = kids
- ❖ app messages = letters in envelopes
- ❖ transport protocol = Ann and Bill who demux to in-house siblings
- ❖ network-layer protocol = postal service
- ❖ Applications: write a book, exchange photos, letters,
- ❖ additional transport services = reliability, encryption
- ❖ another protocol = older kids demux to siblings

# Internet transport-layer protocols

- ❖ reliable, in-order delivery (TCP)
  - congestion control
  - flow control
  - connection setup
- ❖ unreliable, unordered delivery: UDP
  - no-frills extension of “best-effort” IP
- ❖ services not available:
  - delay guarantees
  - bandwidth guarantees



# Chapter 3 outline

3.1 transport-layer services

3.2 multiplexing and demultiplexing

3.3 connectionless transport: UDP

3.4 principles of reliable data transfer

3.5 connection-oriented transport: TCP

- segment structure
- reliable data transfer
- flow control
- connection management

3.6 principles of congestion control

3.7 TCP congestion control

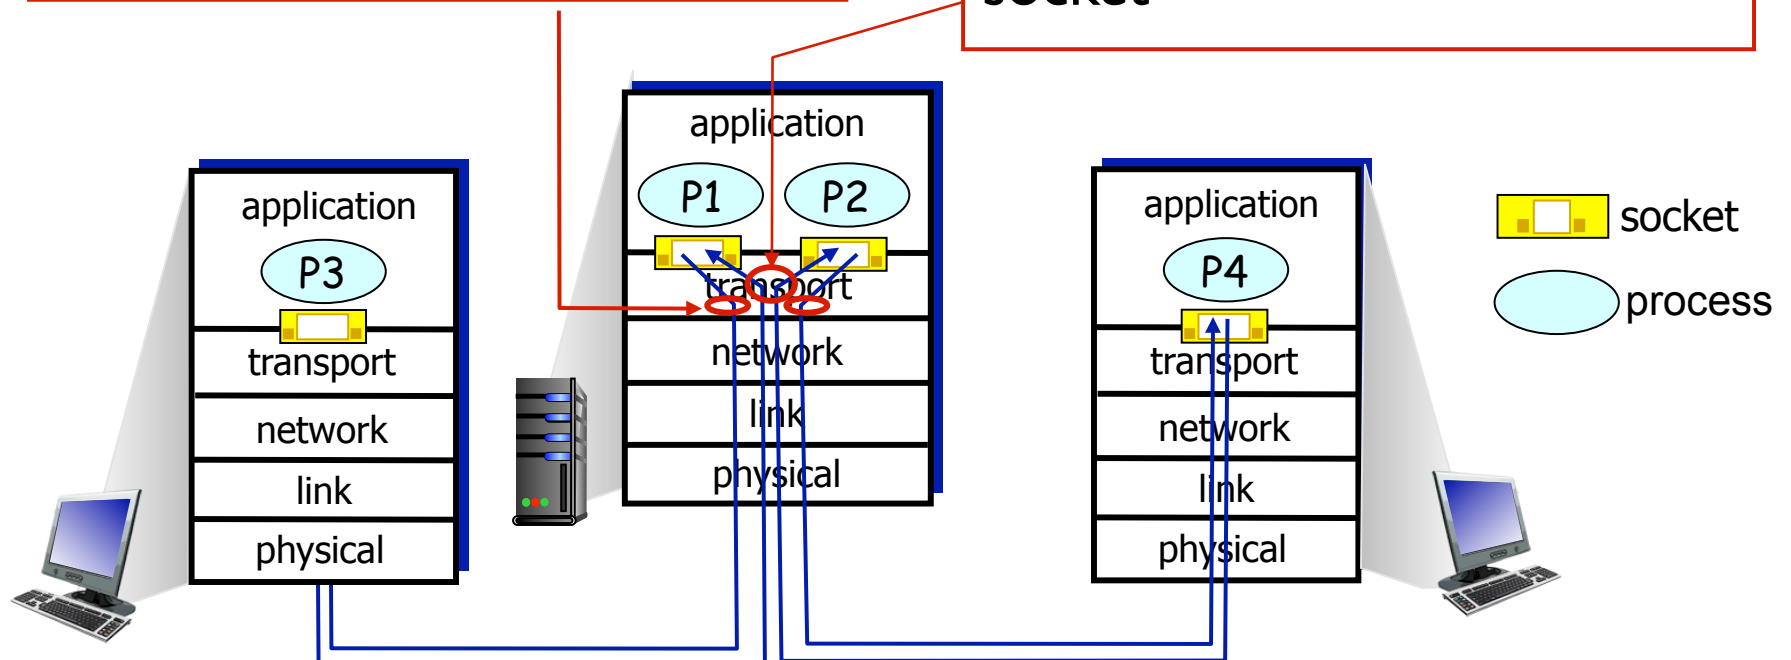
# Multiplexing/demultiplexing

## *multiplexing at sender:*

handle data from multiple sockets, add transport header (later used for demultiplexing)

## *demultiplexing at receiver:*

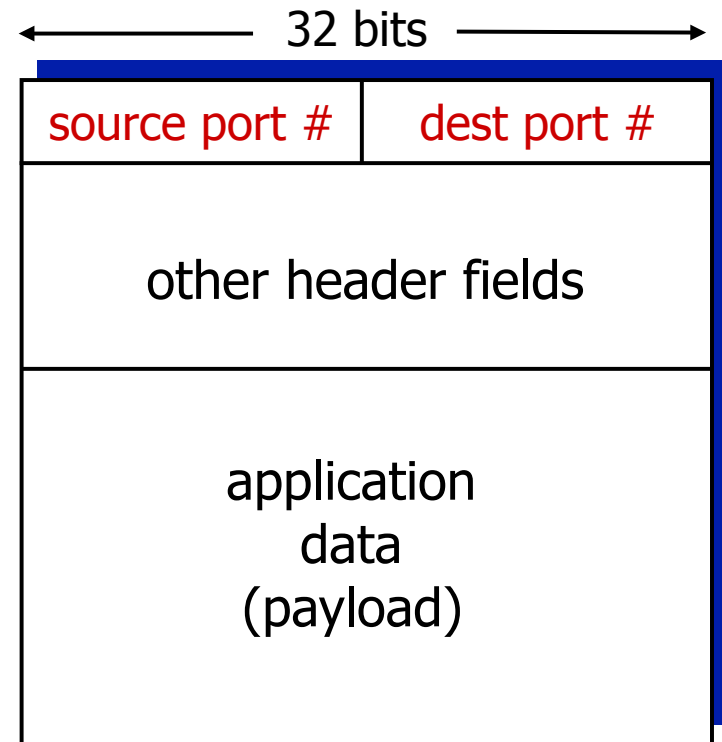
use header info to deliver received segments to correct socket





# TCP/UDP Segment

- ❖ host receives IP datagrams
  - each datagram has source IP address, destination IP address
  - each datagram carries one transport-layer segment
  - each segment has source, destination port number
- ❖ host uses *IP addresses & port numbers* to direct segment to appropriate socket
- ❖ Encapsulation....



TCP/UDP segment format

# Connectionless demultiplexing

- ❖ *recall*: created socket has host-local port #:
    - `DatagramSocket mySocket1 = new DatagramSocket(12534);`
  - ❖ *recall*: when creating datagram to send into UDP socket, must specify dest IP address and destination port #
    - ❖ `mySocket1.sendto(message, (serverName, serverPort))`
- 

- ❖ when host receives UDP segment:

- checks destination port # in segment
- directs UDP segment to socket with that port #



IP datagrams with *same dest. port #*, but different source IP addresses and/or source port numbers will be directed to *same socket* at dest

# Client/server UDP sockets- revisited

## server (running on serverIP)

create socket, port= x:

`serverSocket =  
socket(AF_INET,SOCK_DGRAM)`

↓  
read datagram from  
`serverSocket`

↓  
write reply to  
`serverSocket`  
specifying  
client address,  
port number

## client

create socket:

`clientSocket =  
socket(AF_INET,SOCK_DGRAM)`

↓  
Create datagram with server IP and  
port=x; send datagram via  
`clientSocket`

↓  
read datagram from  
`clientSocket`

↓  
close  
`clientSocket`

# Example app: UDP client

## *Python UDPClient*

include Python's socket  
library →

```
from socket import *  
serverName = 'hostname'  
serverPort = 12000
```

create UDP socket for  
server →

```
clientSocket = socket(socket.AF_INET,  
                      socket.SOCK_DGRAM)
```

get user keyboard  
input →

```
message = raw_input('Input lowercase sentence:')  
clientSocket.sendto(message,(serverName, serverPort))
```

Attach server name, port to  
message; send into socket →

```
modifiedMessage, serverAddress =
```

read reply characters from  
socket into string →

```
clientSocket.recvfrom(2048)
```

```
print modifiedMessage
```

print out received string  
and close socket →

```
clientSocket.close()
```

# Example app: UDP server

## *Python UDPServer*

```
from socket import *
serverPort = 12000

create UDP socket → serverSocket = socket(AF_INET, SOCK_DGRAM)
bind socket to local port
number 12000 → serverSocket.bind(("", serverPort))
print "The server is ready to receive"

while 1:
    loop forever → message, clientAddress = serverSocket.recvfrom(2048)
    Read from UDP socket into
    message, getting client's
    address (client IP and port) → modifiedMessage = message.upper()
    serverSocket.sendto(modifiedMessage, clientAddress)

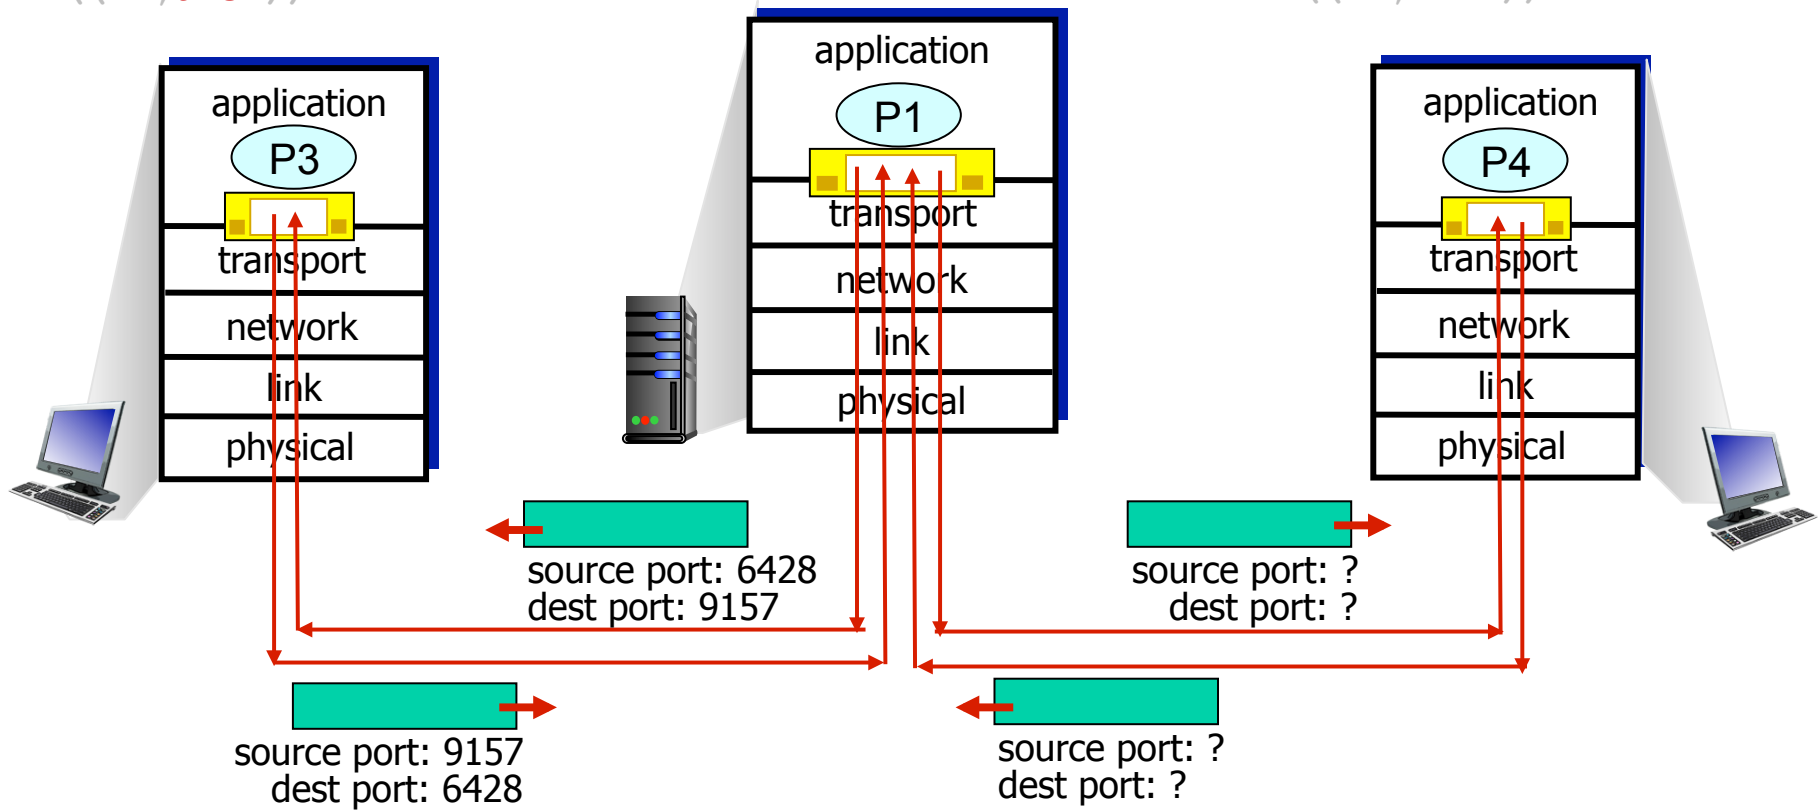
    send upper case string
    back to this client →
```

# Connectionless demux: example

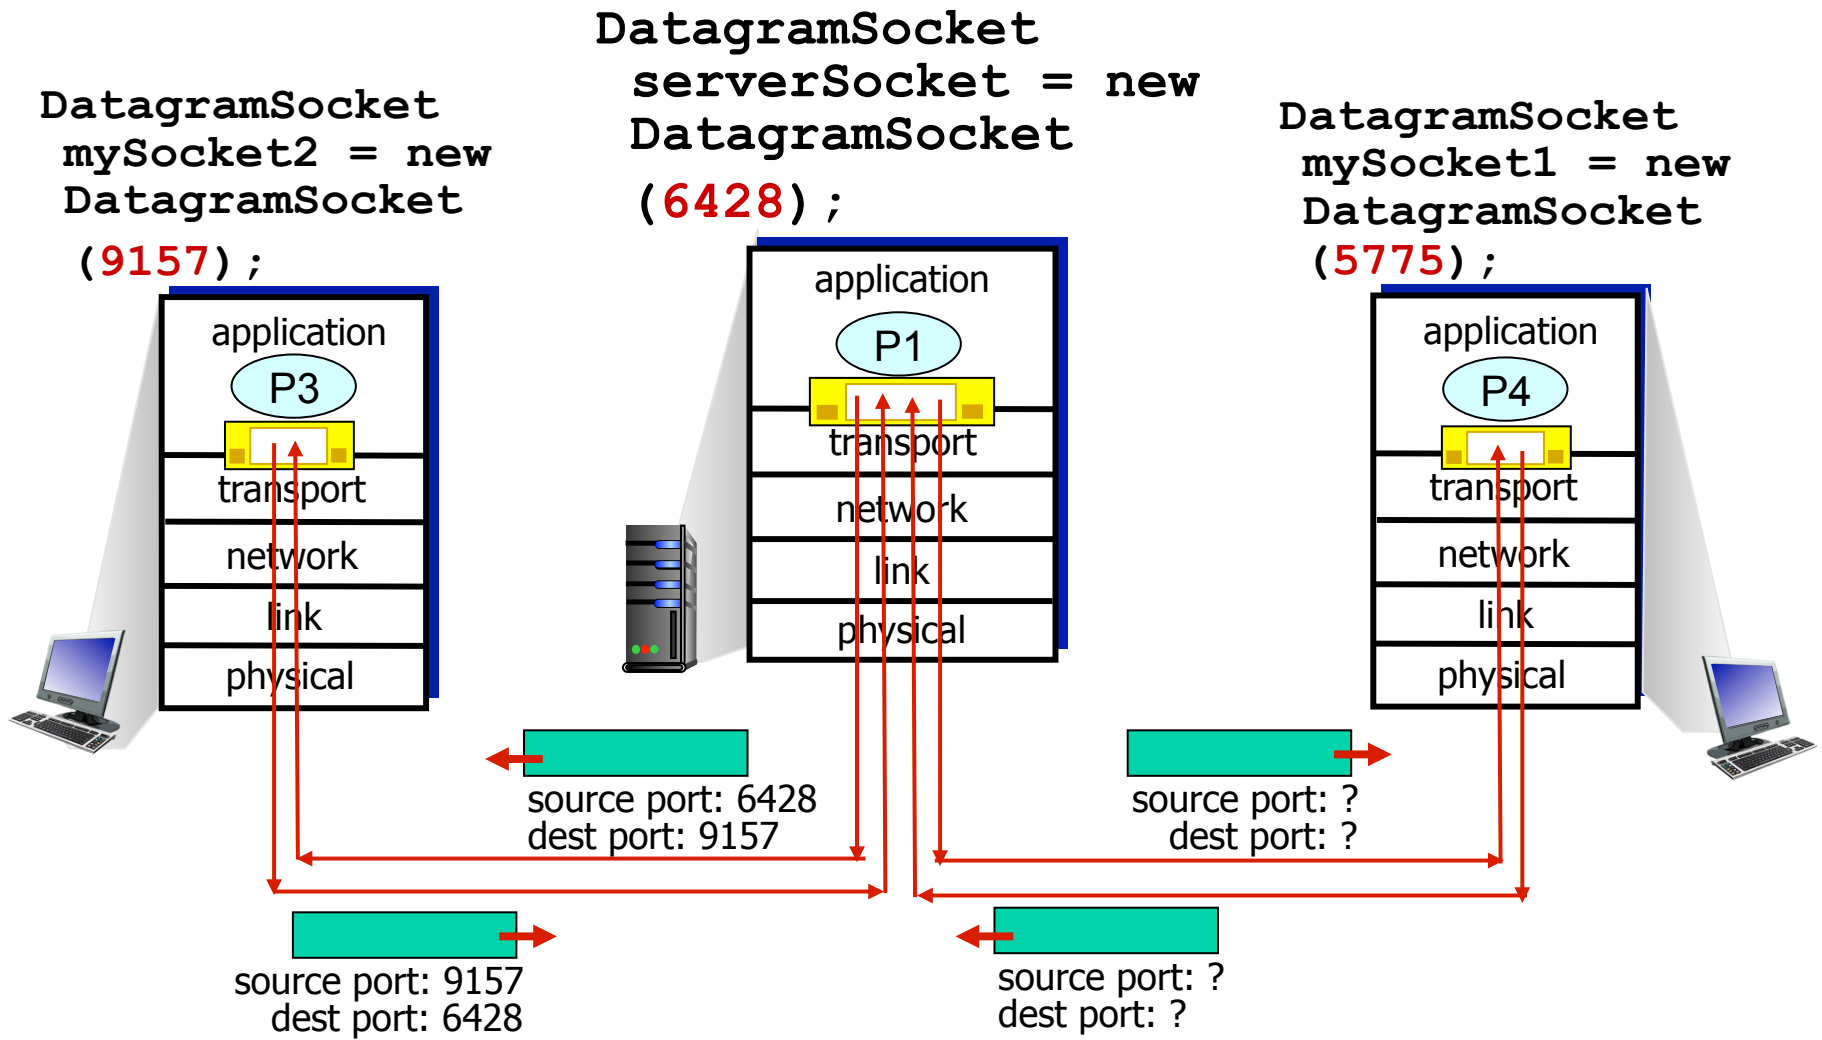
```
clientUDPSocket2 =  
    socket(socket.AF_INET,  
           socket.SOCK_DGRAM)  
clientUDPSocket2.bind  
    ((' ', 9157))
```

```
serverUDPSocket=  
    socket(AF_INET, SOCK_DGRAM)  
serverUDPSocket.bind ((' ',  
                       6428))
```

```
clientUDPSocket2 =  
    socket(socket.AF_INET,  
           socket.SOCK_DGRAM)  
clientUDPSocket2.bind  
    ((' ', 5775))
```



# Connectionless demux: Java example



# Connection-oriented demux

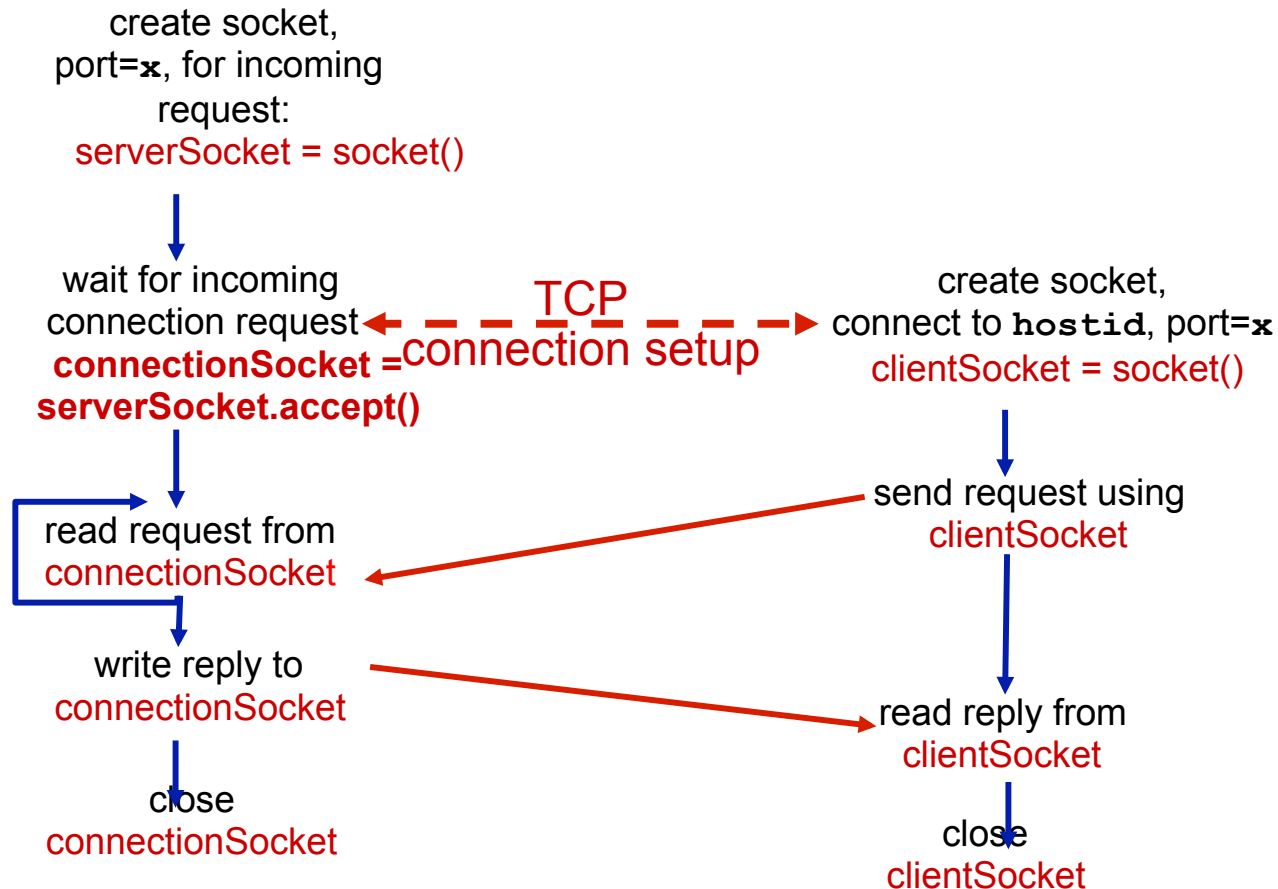
- ❖ TCP socket identified by 4-tuple:
  - source IP address
  - source port number
  - dest IP address
  - dest port number
- ❖ demux: receiver uses all four values to direct segment to appropriate socket
- ❖ server host may support many simultaneous TCP sockets:
  - each socket identified by its own 4-tuple
- ❖ web servers have different sockets for each connecting client
  - non-persistent HTTP will have different socket for each request



# Client/server TCP sockets – revisited

server (running on `hostid`)

client



# Example app:TCP client

## *Python TCPClient*

```
from socket import *
serverName = 'servername'
serverPort = 12000
clientSocket = socket(AF_INET, SOCK_STREAM)
clientSocket.connect((serverName,serverPort))
sentence = raw_input('Input lowercase sentence:')
clientSocket.send(sentence)
modifiedSentence = clientSocket.recv(1024)
print 'From Server:', modifiedSentence
clientSocket.close()
```

create TCP socket for  
server, remote port 12000 →

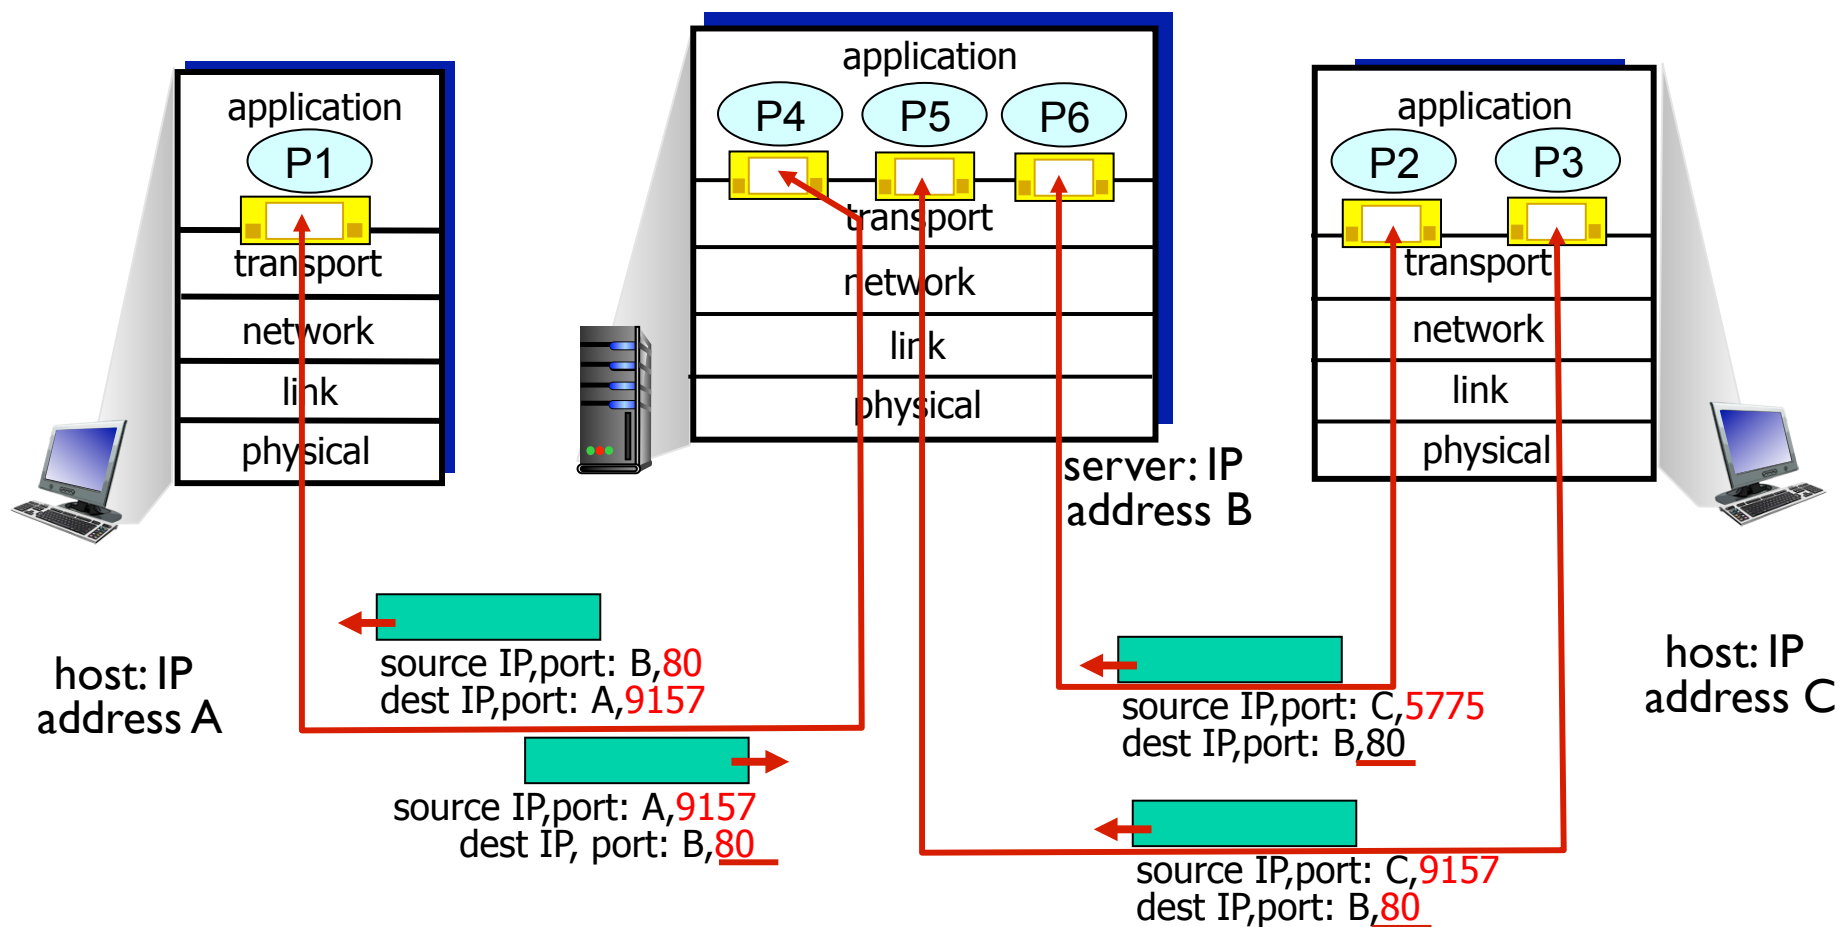
No need to attach server  
name, port →

# Example app: TCP server

## *Python TCPServer*

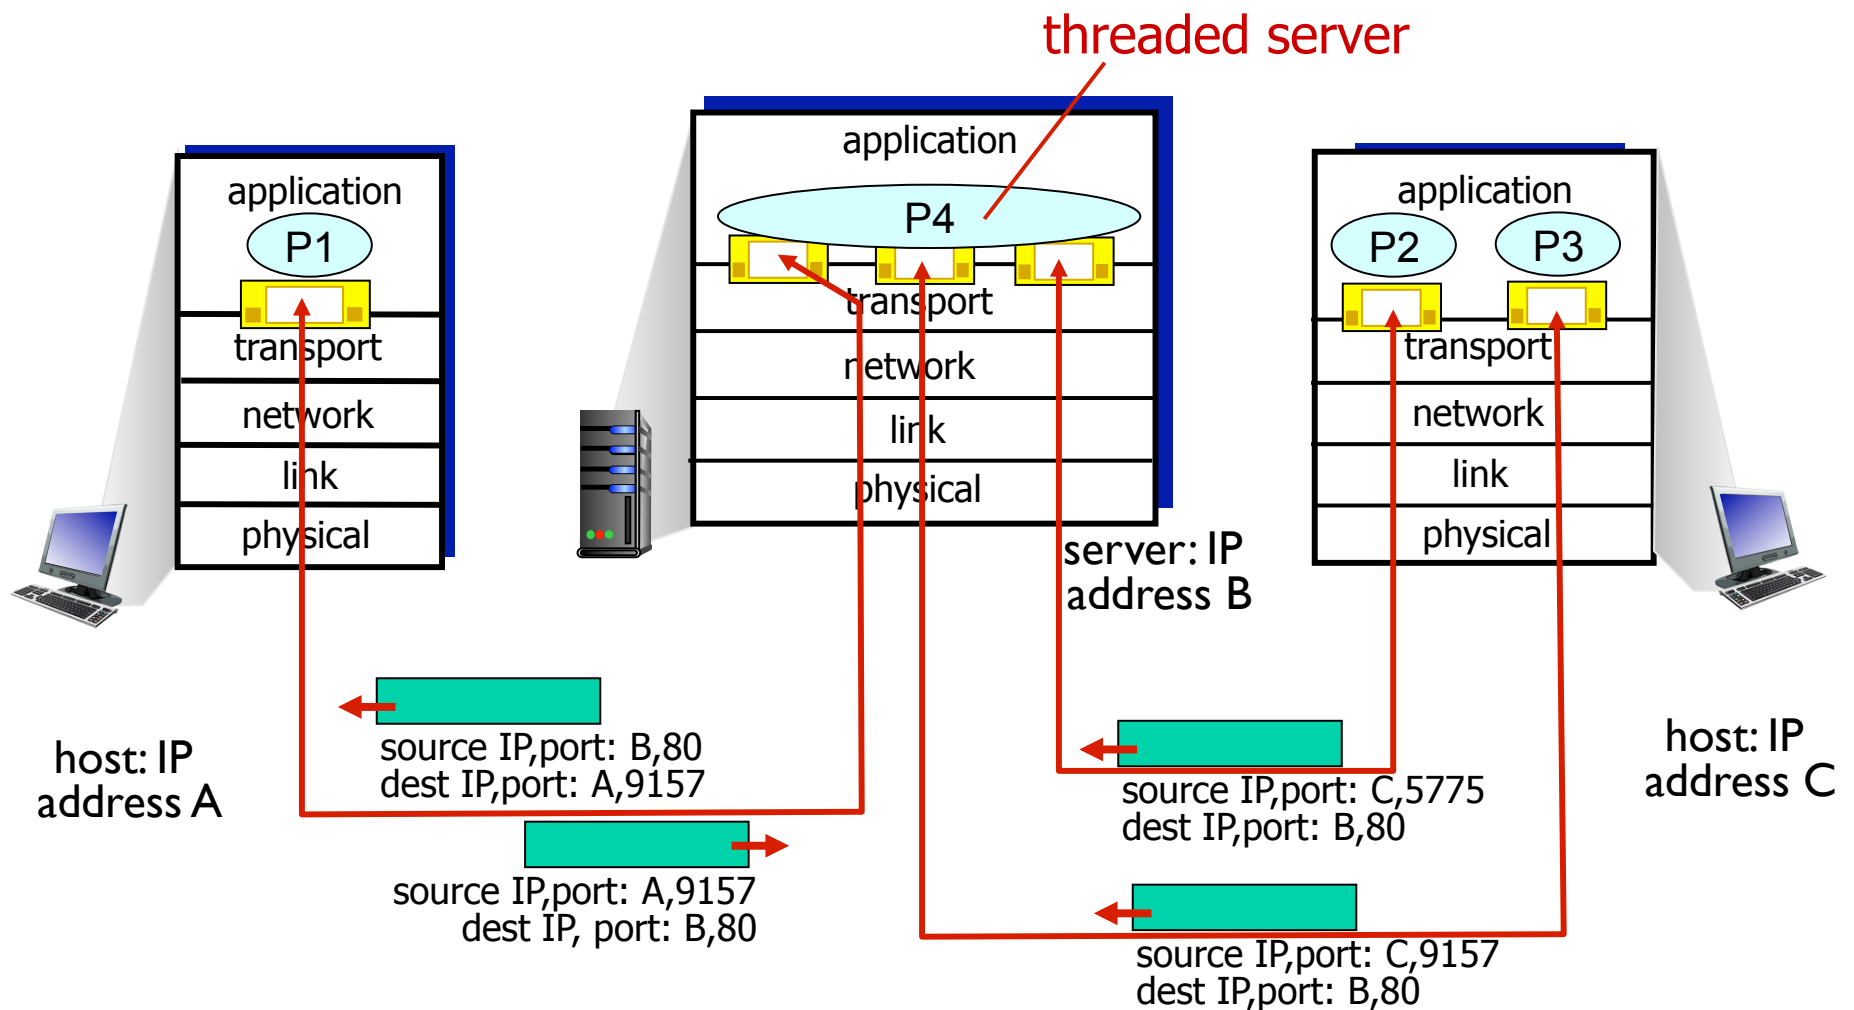
create TCP welcoming socket	→	<pre>from socket import * serverPort = 12000 serverSocket = socket(AF_INET, SOCK_STREAM) serverSocket.bind(('', serverPort)) serverSocket.listen(1) print 'The server is ready to receive'while 1:</pre>
server begins listening for incoming TCP requests	→	<pre>connectionSocket, addr = serverSocket.accept()</pre>
loop forever	→	<pre>    sentence = connectionSocket.recv(1024)     capitalizedSentence = sentence.upper()     connectionSocket.send(capitalizedSentence)     connectionSocket.close()</pre>
server waits on accept() for incoming requests, new socket created on return	→	
read bytes from socket (but not address as in UDP)	→	
close connection to this client (but <i>not</i> welcoming socket)	→	

# Connection-oriented demux: example



three segments, all destined to IP address: B,  
dest port: 80 are demultiplexed to *different* sockets

# Connection-oriented demux: example



# Chapter 3 outline

3.1 transport-layer services

3.2 multiplexing and demultiplexing

3.3 connectionless transport: UDP

3.4 principles of reliable data transfer

3.5 connection-oriented transport: TCP

- segment structure
- reliable data transfer
- flow control
- connection management

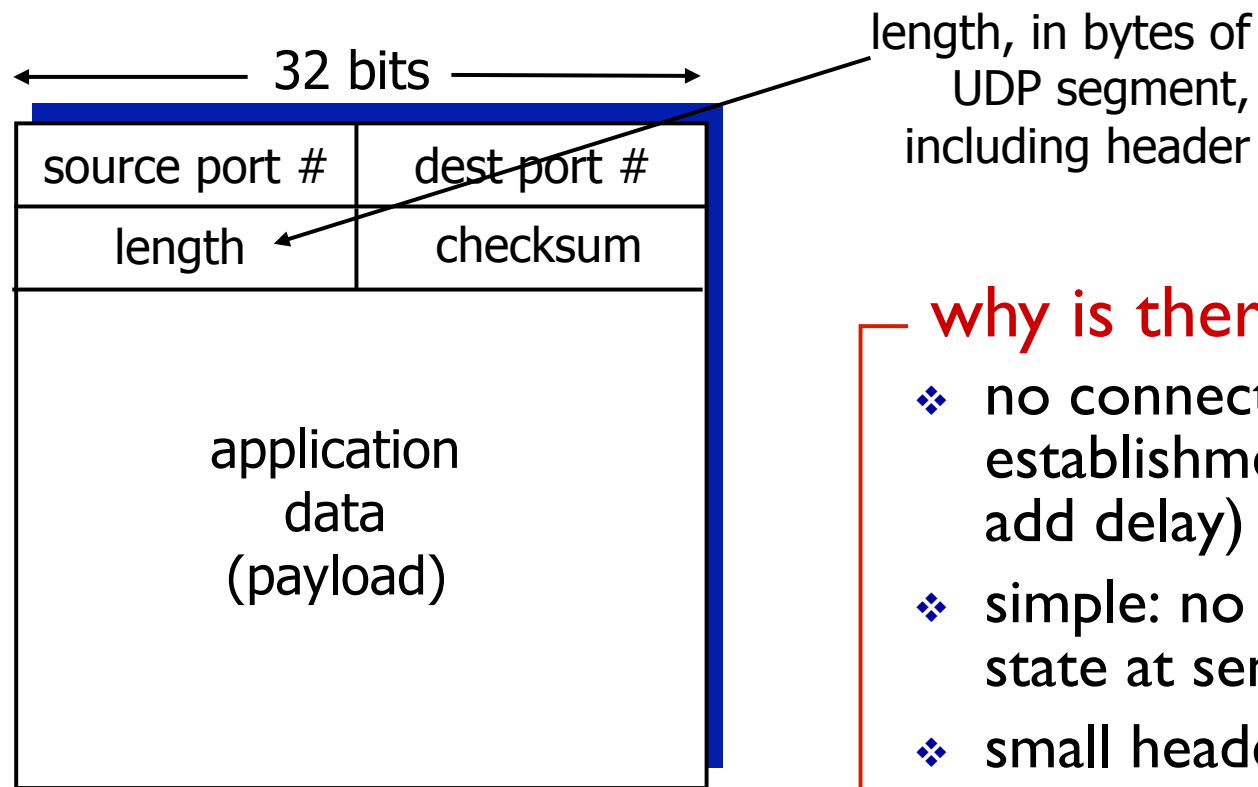
3.6 principles of congestion control

3.7 TCP congestion control

# UDP: User Datagram Protocol [RFC 768]

- ❖ “no frills,” “bare bones”  
Internet transport  
protocol
- ❖ “best effort” service,  
UDP segments may be:
  - lost
  - delivered out-of-order  
to app
- ❖ *connectionless*:
  - no handshaking  
between UDP sender,  
receiver
  - each UDP segment  
handled independently  
of others
- ❖ UDP use:
  - streaming multimedia  
apps (loss tolerant, rate  
sensitive)
  - DNS
  - SNMP
- ❖ reliable transfer over  
UDP:
  - add reliability at  
application layer
  - application-specific error  
recovery!

# UDP: segment header (8B)



UDP segment format

## why is there a UDP?

- ❖ no connection establishment (which can add delay)
- ❖ simple: no connection state at sender, receiver
- ❖ small header size
- ❖ no congestion control: UDP can blast away as fast as desired



# UDP checksum

*Goal:* detect “errors” (e.g., flipped bits) in transmitted segment

## sender:

- ❖ treat segment contents, including header fields, as sequence of 16-bit integers
- ❖ checksum: addition (one's complement sum) of segment contents
- ❖ sender puts checksum value into UDP checksum field

## receiver:

- ❖ compute checksum of received segment
- ❖ check if computed checksum equals checksum field value:
  - NO - error detected
  - YES - no error detected.  
*But maybe errors nonetheless? More later*  
....

# Internet checksum: example

example: add two 16-bit integers

	1	1	1	0	0	1	1	0	0	1	1	0	0	1	1	0
	1	1	0	1	0	1	0	1	0	1	0	1	0	1	0	1
<hr/>																
wraparound	1	1	0	1	1	1	0	1	1	1	0	1	1	1	0	1
<hr/>																
sum	1	0	1	1	1	0	1	1	1	0	1	1	1	1	0	0
checksum	0	1	0	0	0	1	0	0	0	1	0	0	0	0	1	1

*Note:* when adding numbers, a carryout from the most significant bit needs to be added to the result