

# Chapter 3 outline

3.1 transport-layer services

3.2 multiplexing and demultiplexing

3.3 connectionless transport: UDP

3.4 principles of reliable data transfer

3.5 connection-oriented transport: TCP


- segment structure
- reliable data transfer
- flow control
- connection management

3.6 principles of congestion control

3.7 TCP congestion control

# Principles of reliable data transfer

- ❖ important in application, transport, link layers
  - top-10 list of important networking topics!

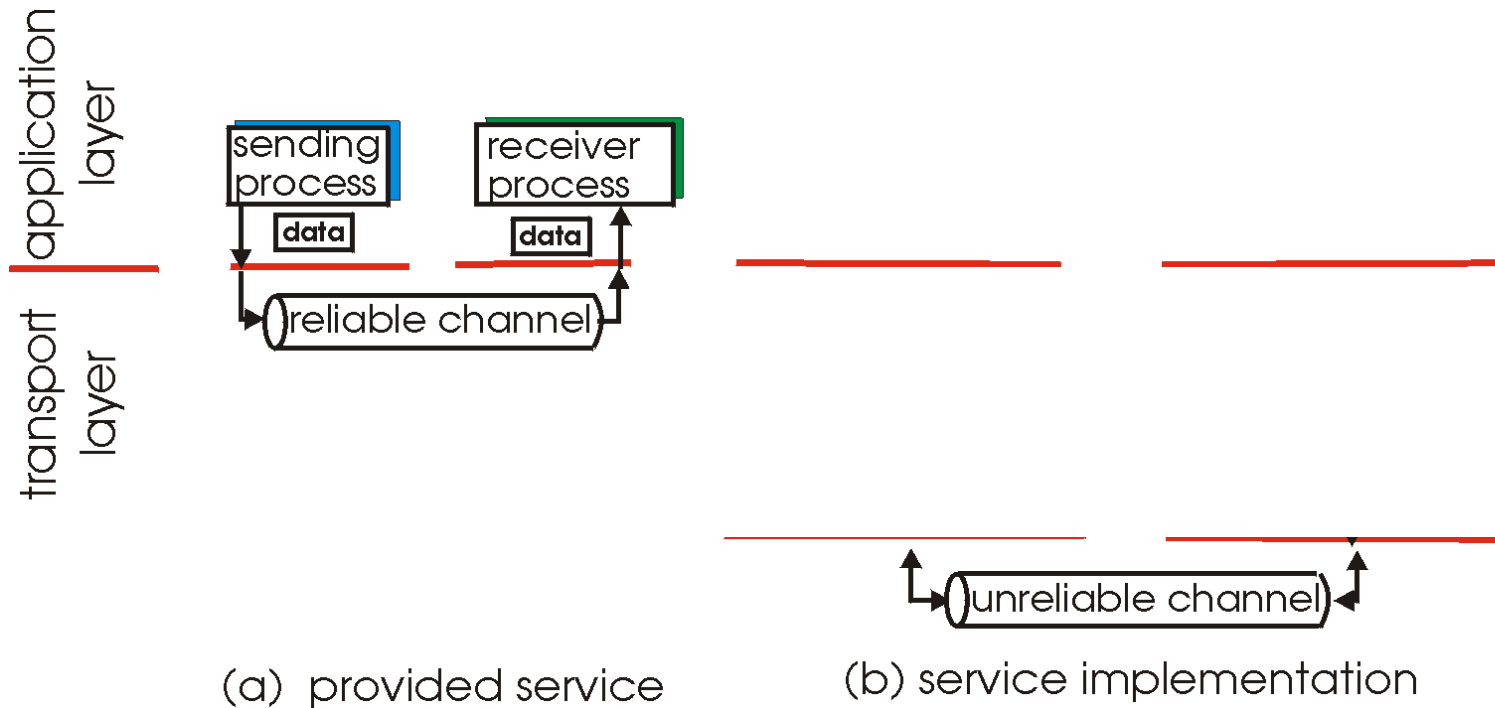


(a) provided service

- ❖ characteristics of unreliable channel will determine complexity of reliable data transfer protocol (rdt)

# Principles of reliable data transfer

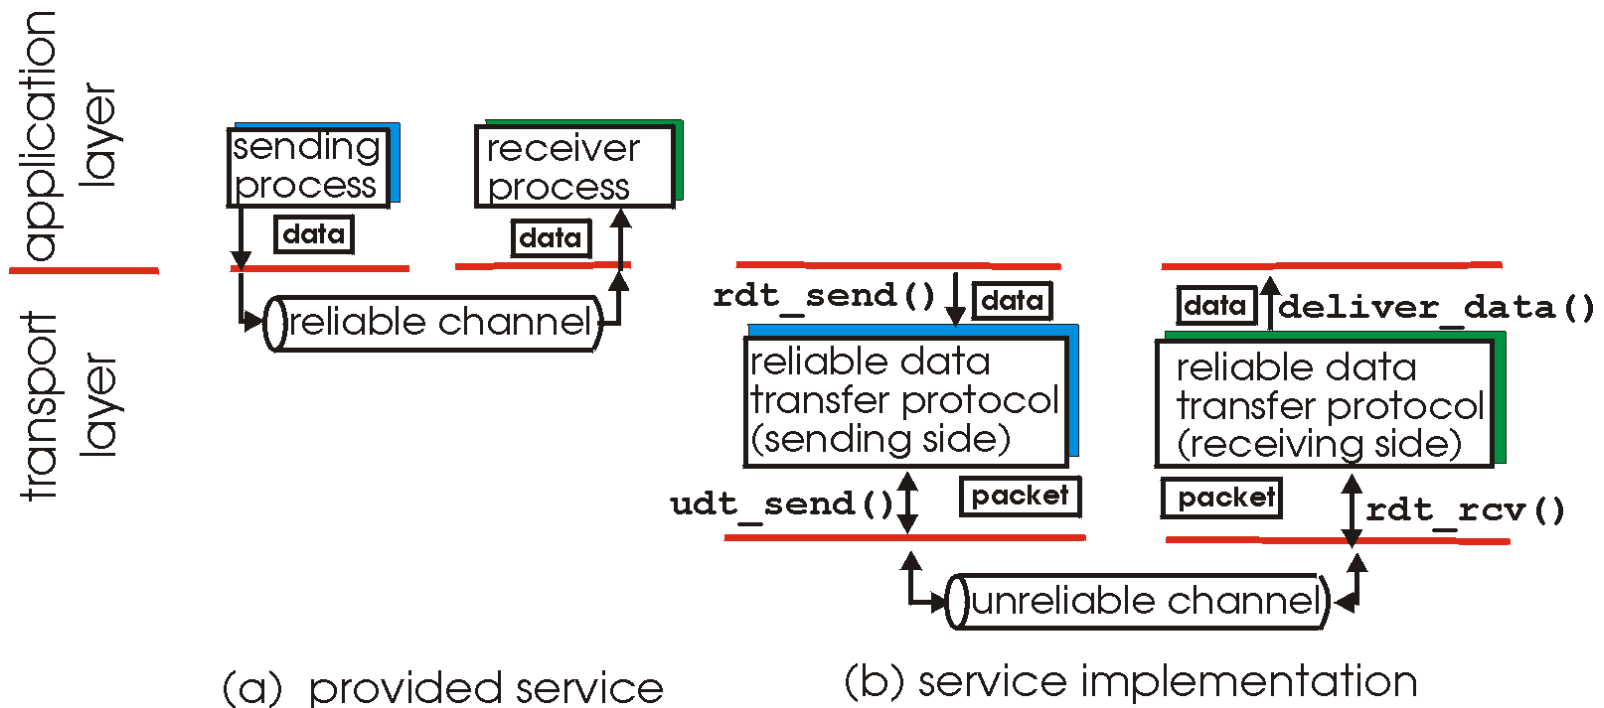
- ❖ important in application, transport, link layers
  - top-10 list of important networking topics!



- ❖ characteristics of unreliable channel will determine complexity of reliable data transfer protocol (rdt)

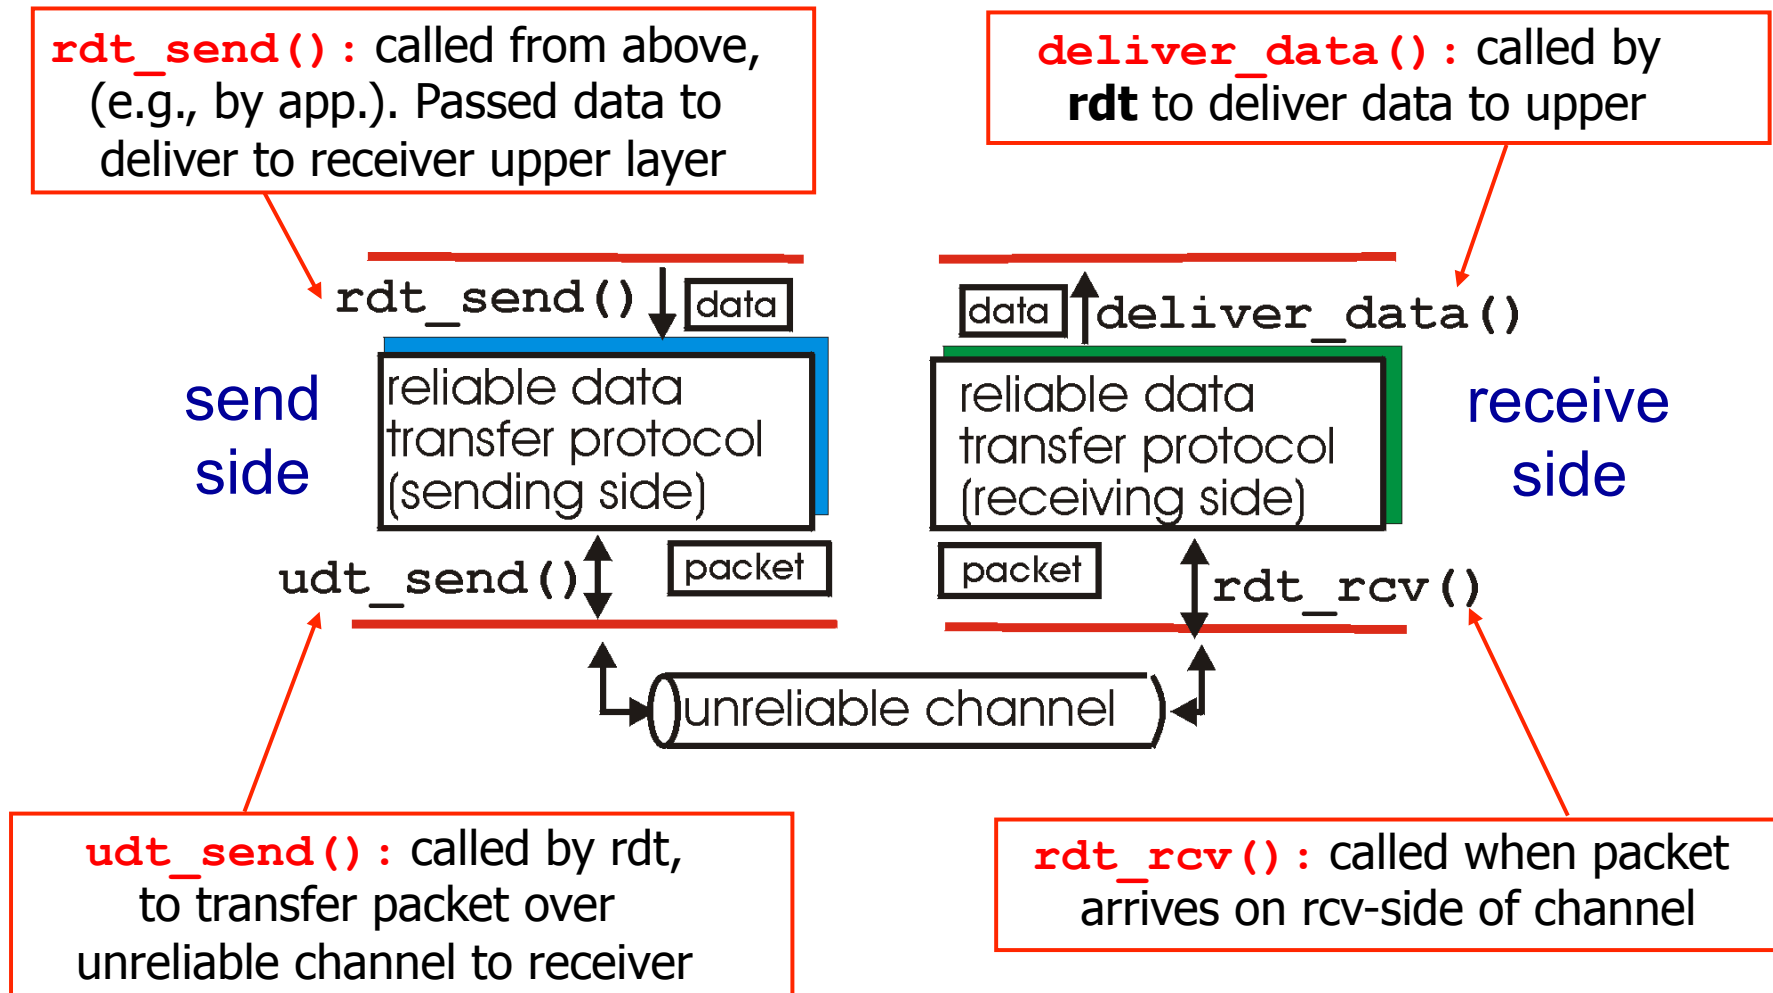
# Principles of reliable data transfer

- ❖ important in application, transport, link layers
  - top-10 list of important networking topics!



- ❖ characteristics of unreliable channel will determine complexity of reliable data transfer protocol (rdt)

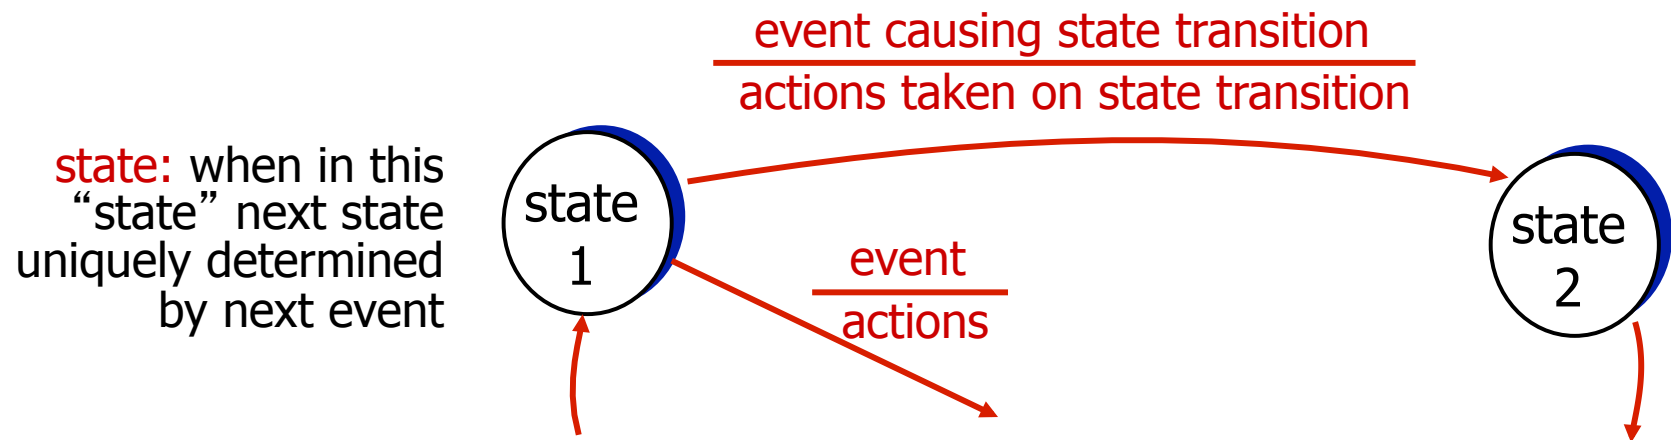
# Reliable data transfer: getting started



# Reliable data transfer: getting started


We will:

- ❖ incrementally develop sender, receiver sides of reliable data transfer protocol (rdt)
- ❖ consider only unidirectional data transfer
  - but control info will flow on both directions!
- ❖ use finite state machines (FSM) to specify sender, receiver




# rdt1.0: reliable transfer over a reliable channel

- ❖ underlying channel perfectly reliable
  - no bit errors
  - no loss of packets
- ❖ separate FSMs for sender, receiver:
  - sender sends data into underlying channel
  - receiver reads data from underlying channel
- ❖ Clearly RDT is useless here 😊



sender



receiver

# rdt2.0: channel with bit errors

- ❖ underlying channel may flip bits in packet
  - checksum to detect bit errors
- ❖ *the question: how to recover from errors:*

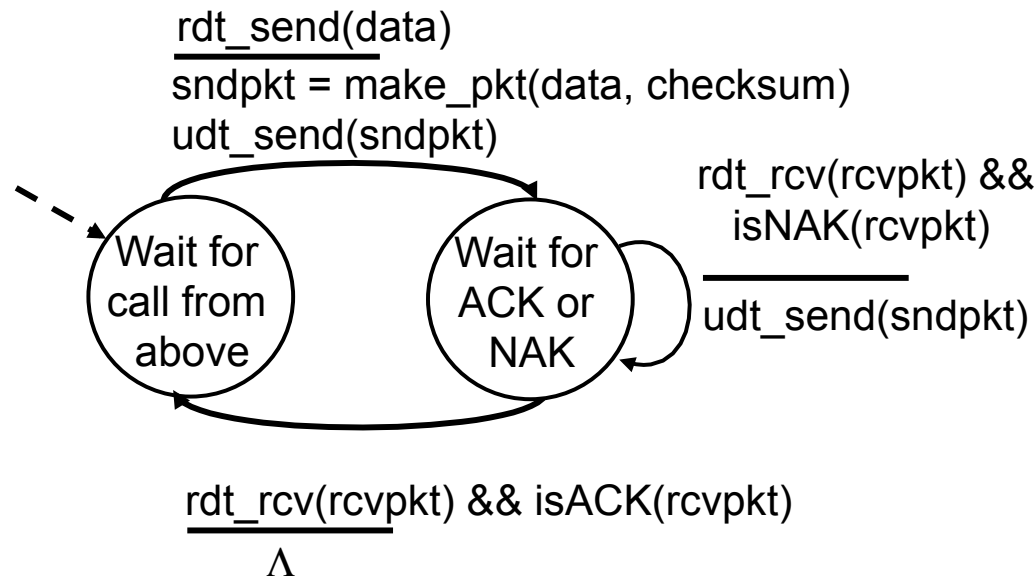
*How do humans recover from “errors”  
during conversation?*



# rdt2.0: channel with bit errors


- ❖ underlying channel may flip bits in packet
  - checksum to detect bit errors
- ❖ Q: how to recover from errors?
- ❖ A: new mechanisms needed:
  - *feedback (control msgs from receiver to sender)*
    - *acknowledgements (ACKs)*: receiver explicitly tells sender that pkt received OK
    - *negative acknowledgements (NAKs)*: receiver explicitly tells sender that pkt had errors
  - *retransmission*
    - sender retransmits pkt on receipt of NAK

# rdt2.0: FSM specification



sender

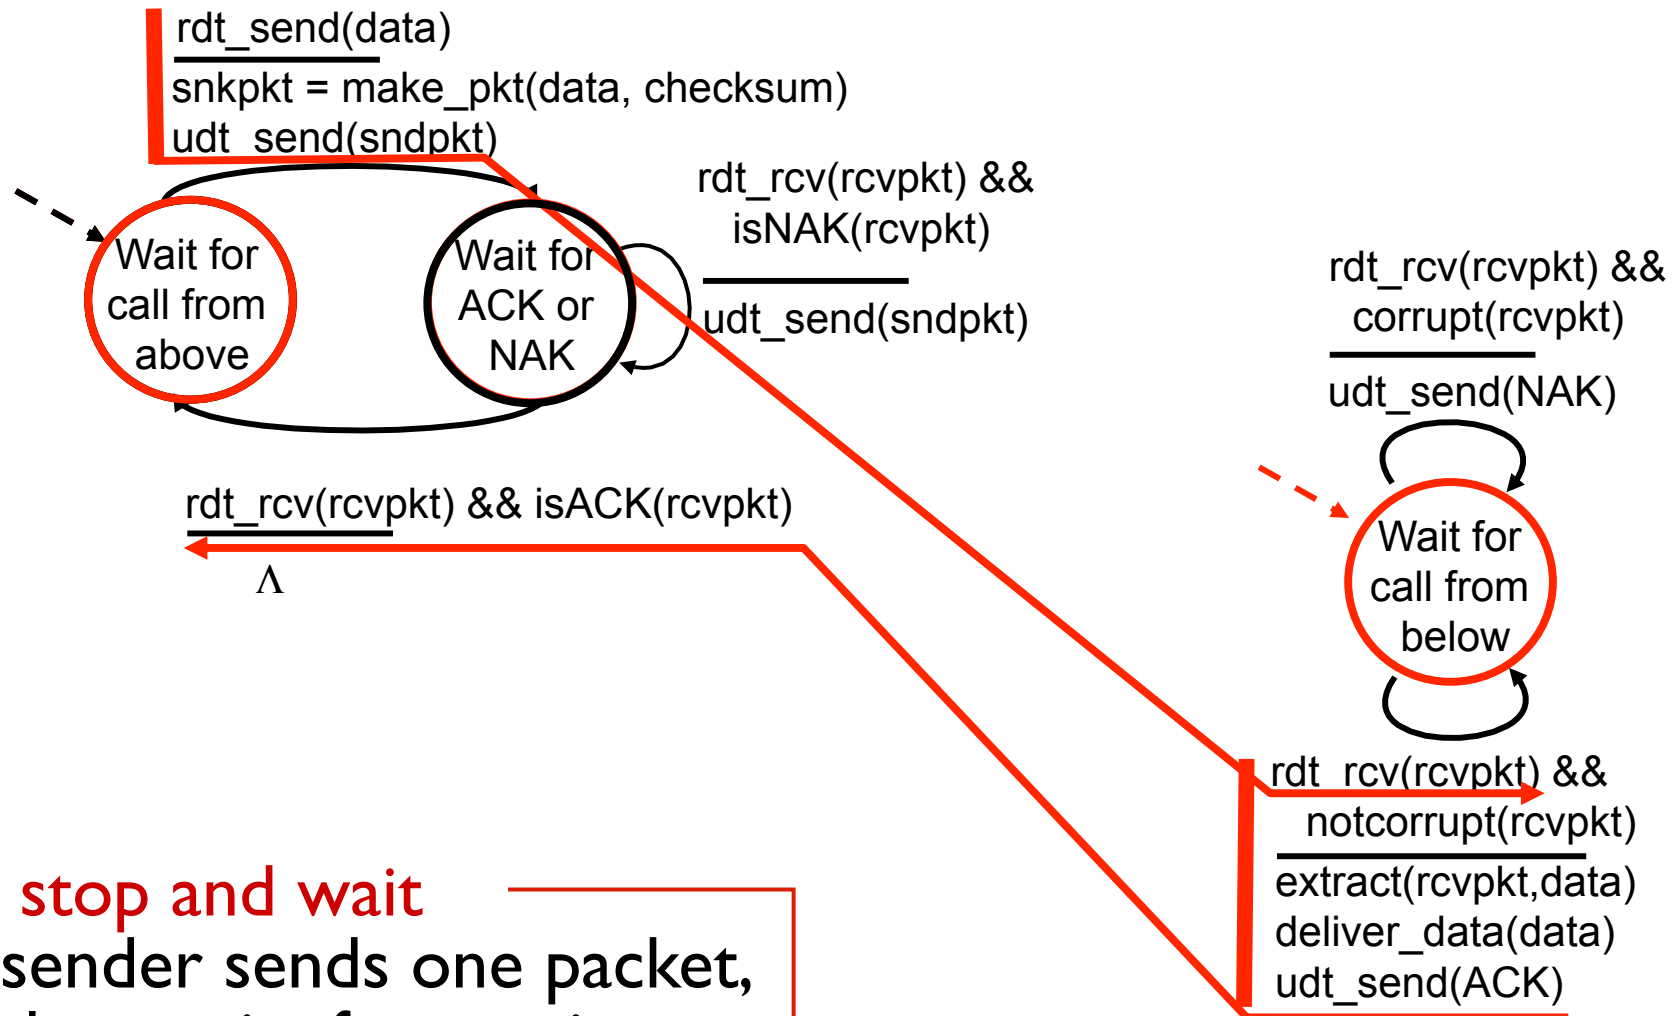
receiver



## stop and wait

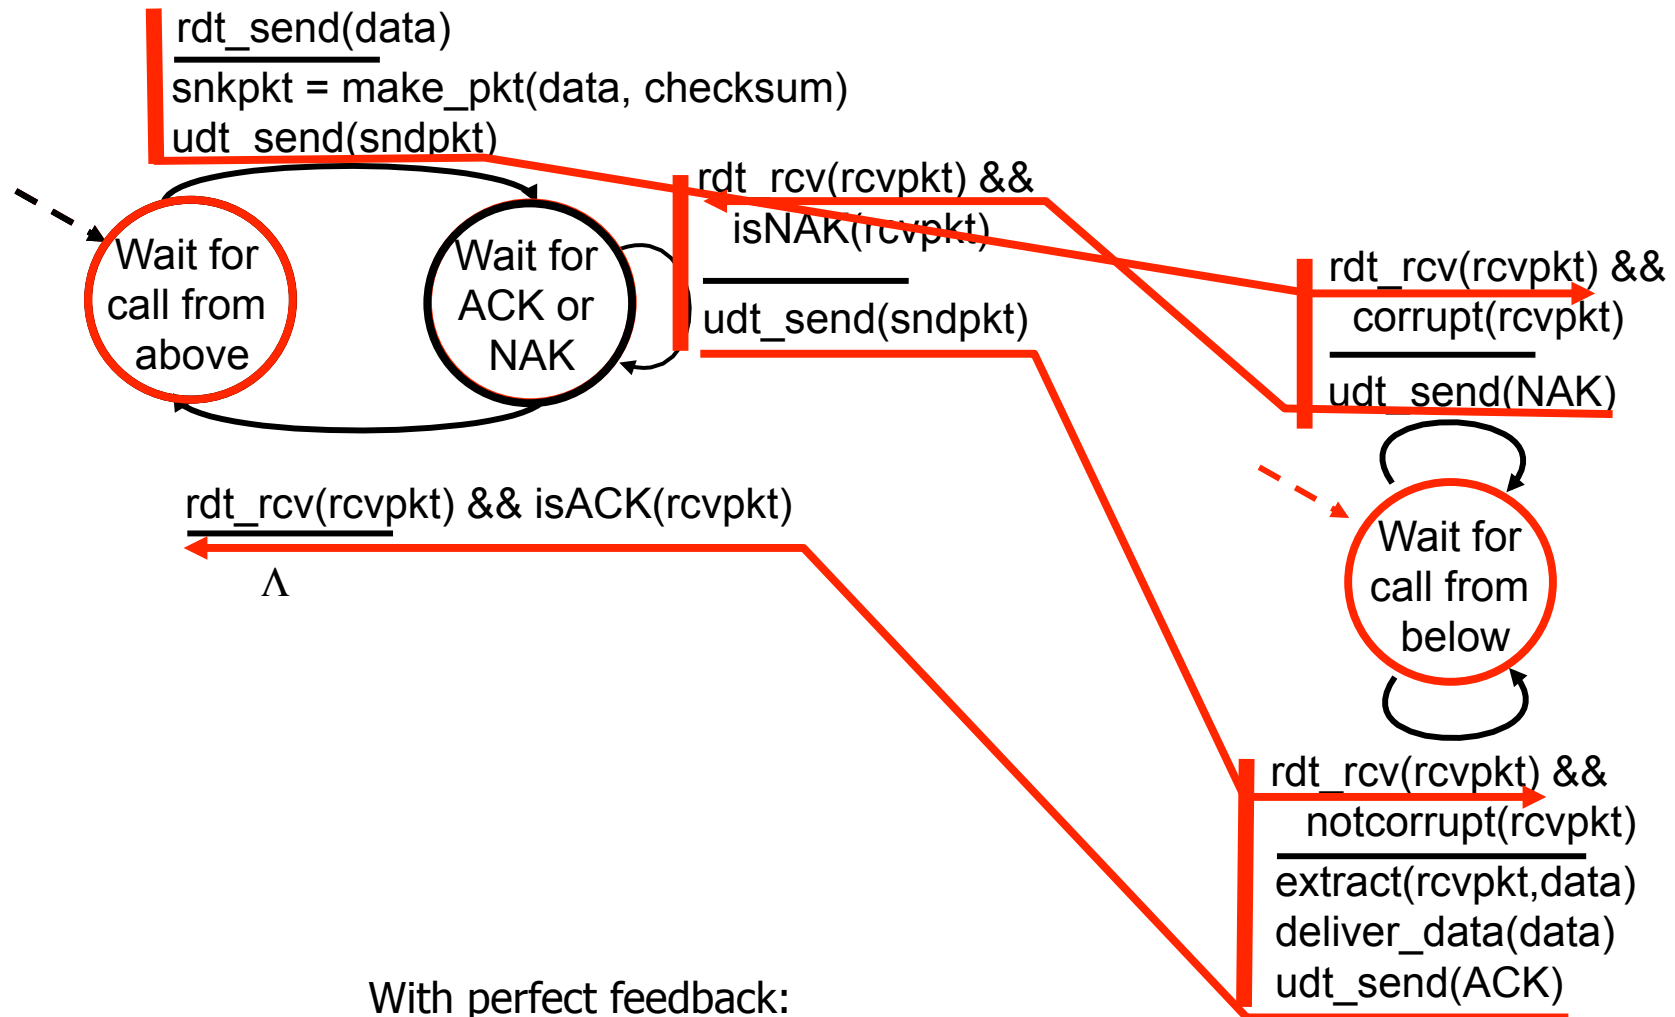
sender sends one packet,  
then waits for receiver  
response

# rdt2.0: operation with no errors



**stop and wait**  
sender sends one packet,  
then waits for receiver  
response

# rdt2.0: error scenario



With perfect feedback:  
sender's state and receiver's state are in perfect sync

# rdt2.0 has a fatal flaw!

## what happens if ACK/NAK corrupted?

- ❖ sender doesn't know what happened at receiver!
  - Packet received or corrupted?
  - ACK corrupted not (exactly) the same as ACK being lost
- ❖ solutions
  - Feedback on feedback...
    - may be corrupted but we can detect that (e.g. through checksum).
  - Error detection+correction
    - Cannot handle lost acks
  - Retransmissions
    - can't just retransmit: possible duplicates

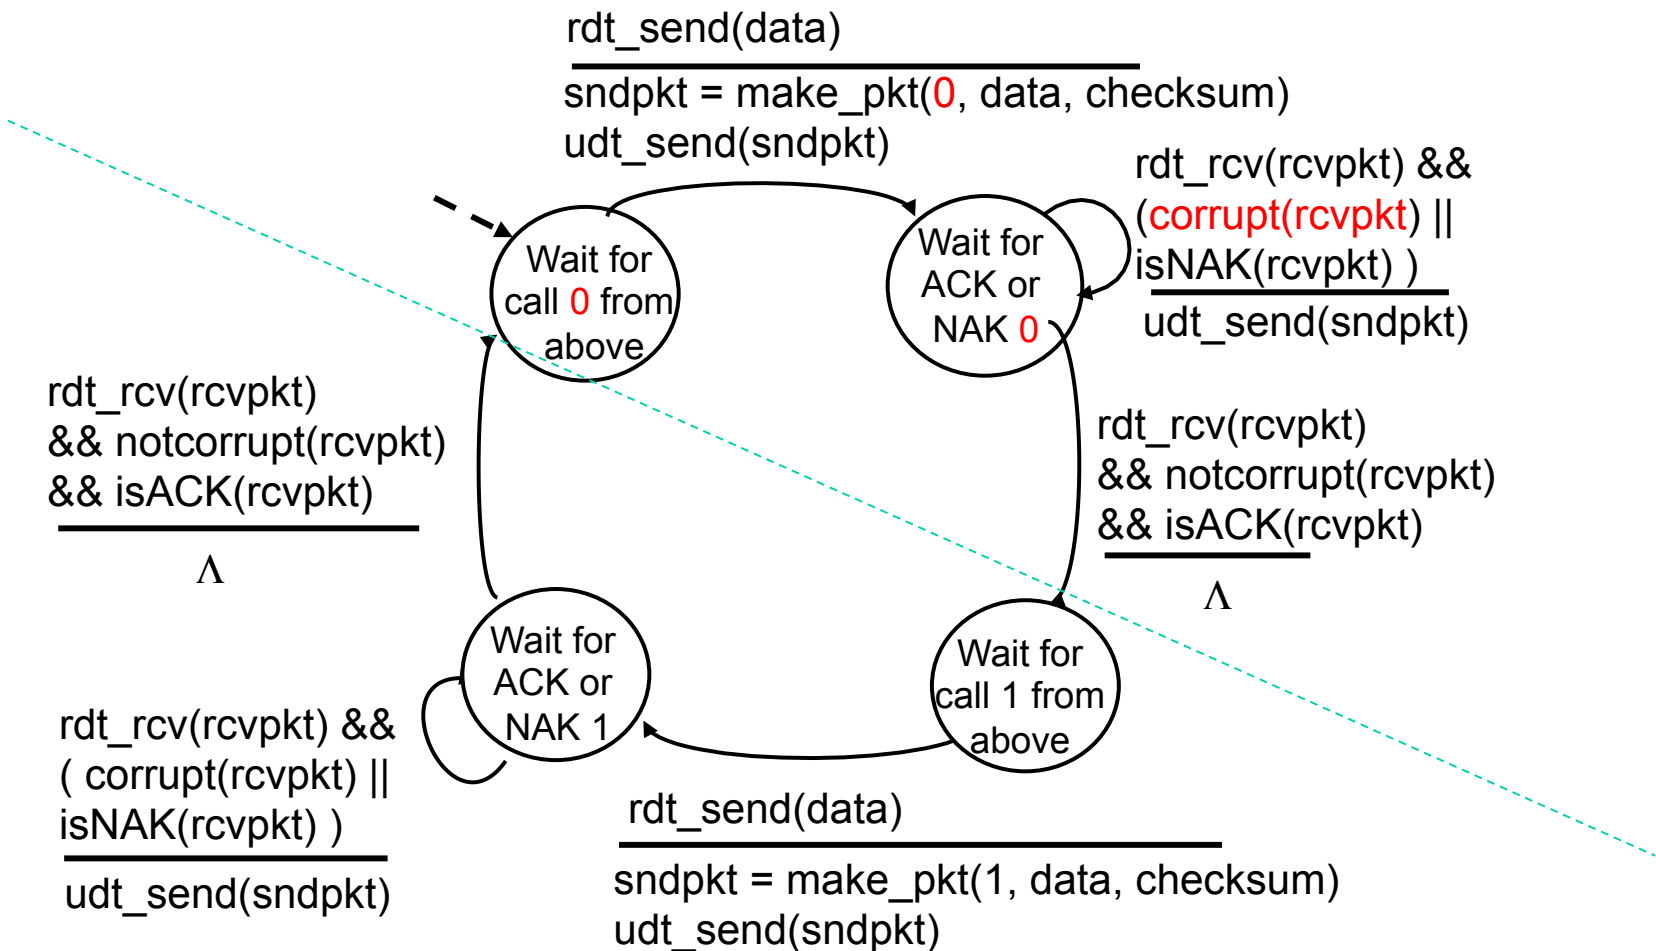
## handling duplicates:

- ❖ sender **conservatively** retransmits current pkt if ACK/NAK corrupted
- ❖ sender adds **sequence number** to each pkt
- ❖ receiver discards (doesn't deliver up) duplicate pkt
- ❖ **but receiver always ACKs/NAKs received packet (to help the server move on to next seqno)**

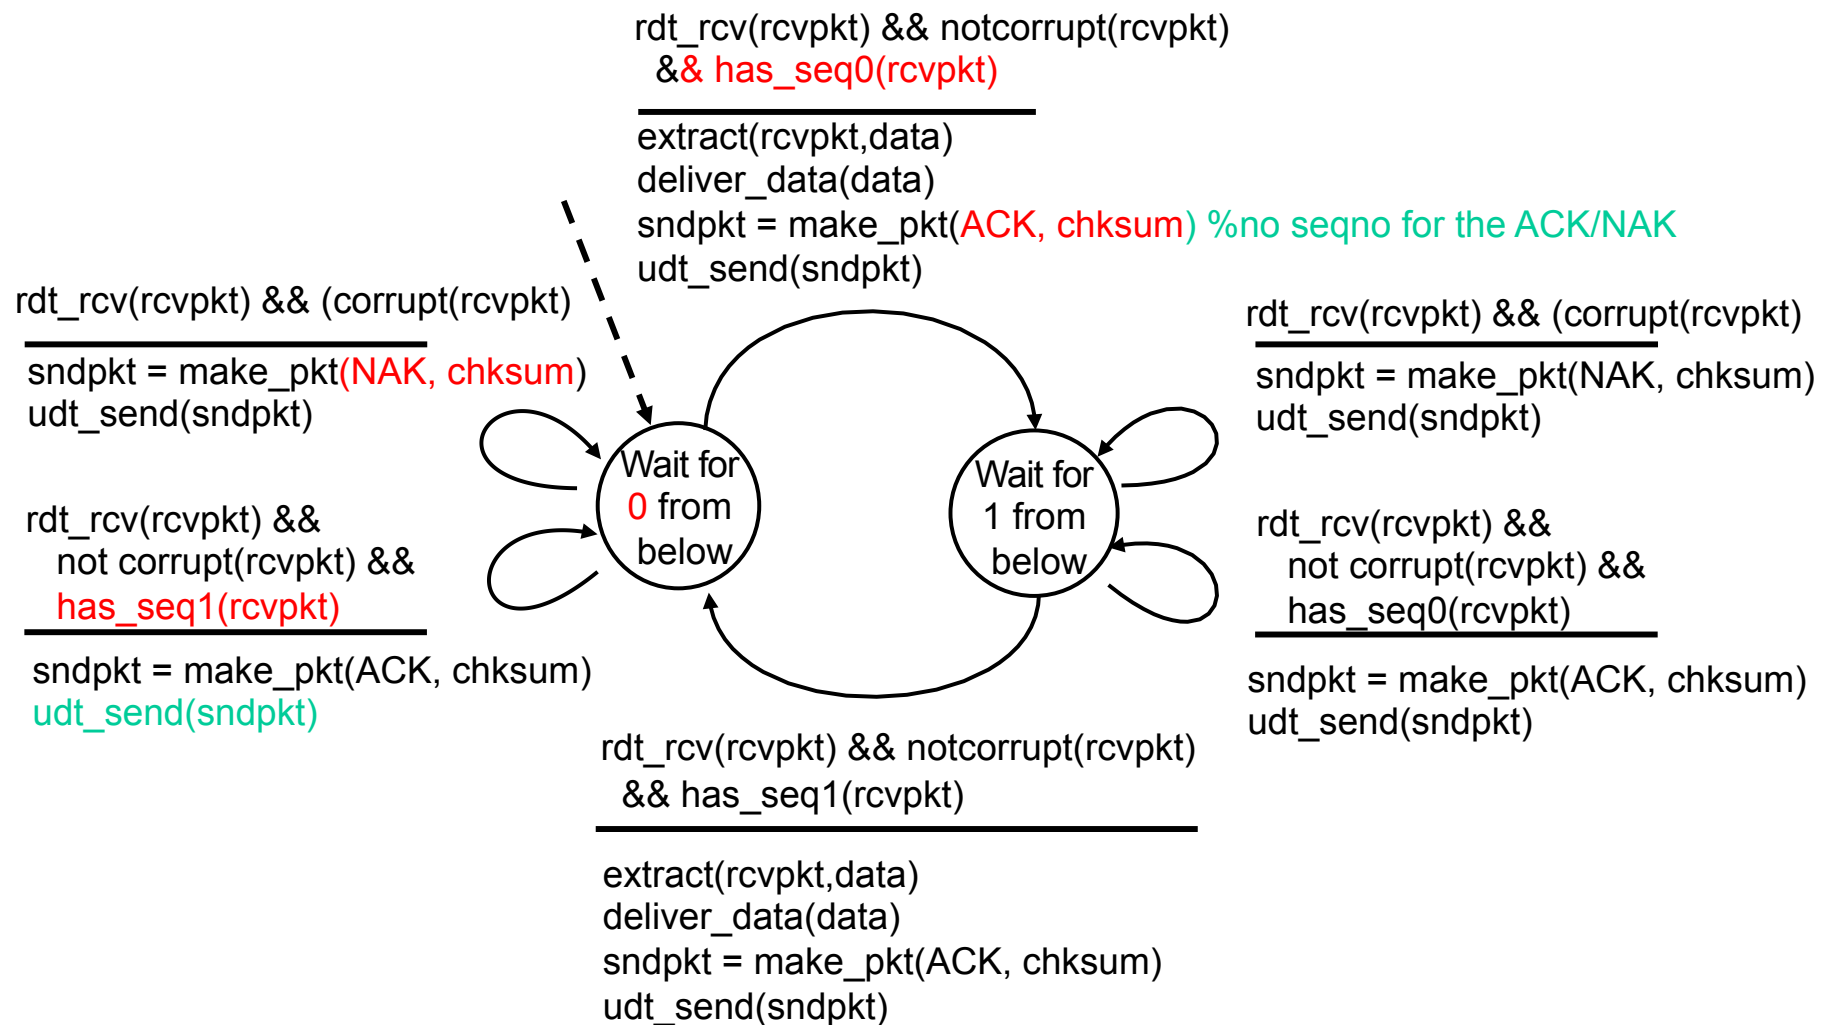
## stop and wait

sender sends one packet, then waits for receiver response

# rdt2.1: sender, handles garbled ACK/NAKs



# rdt2.1: receiver, handles garbled ACK/NAKs



# rdt2.1: discussion

## sender:


- ❖ must check if received ACK/NAK corrupted
- ❖ seq # added to pkt
  - two seq. #'s (0,1) will suffice for stop&wait
  - state indicates whether sender retransmits previous or transmits new pkt
  - twice as many states

## receiver:

- ❖ must check if received packet is duplicate
  - state indicates whether 0 or 1 is expected pkt seq #
  - note: receiver can *not* know if its last ACK/NAK received OK at sender
- ❖ ACKs/NAKs don't need seq# for channel that don't drop packets
  - Always refer to most recently sent pkt
  - Seqnos necessary in ACK-only protocols and when channel drops packets

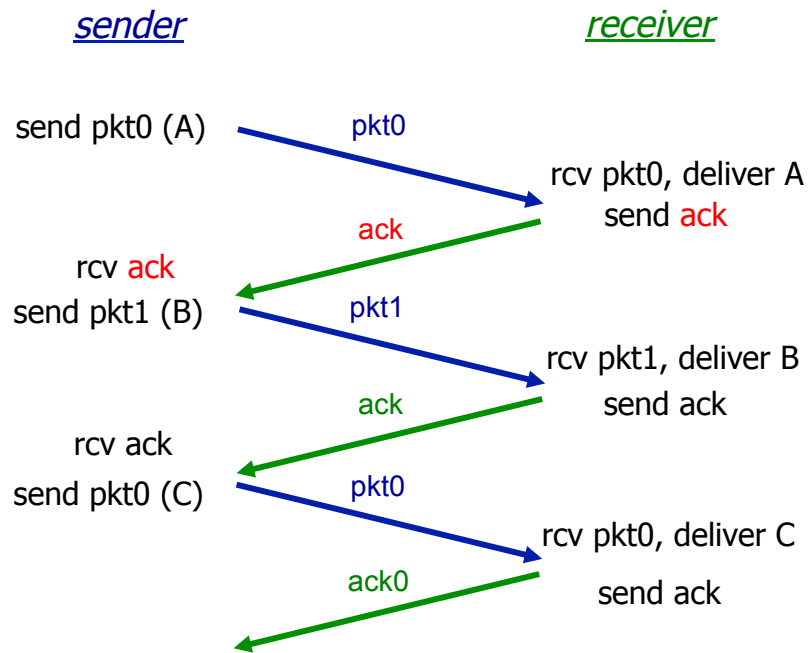


# Rdt2.1 in action

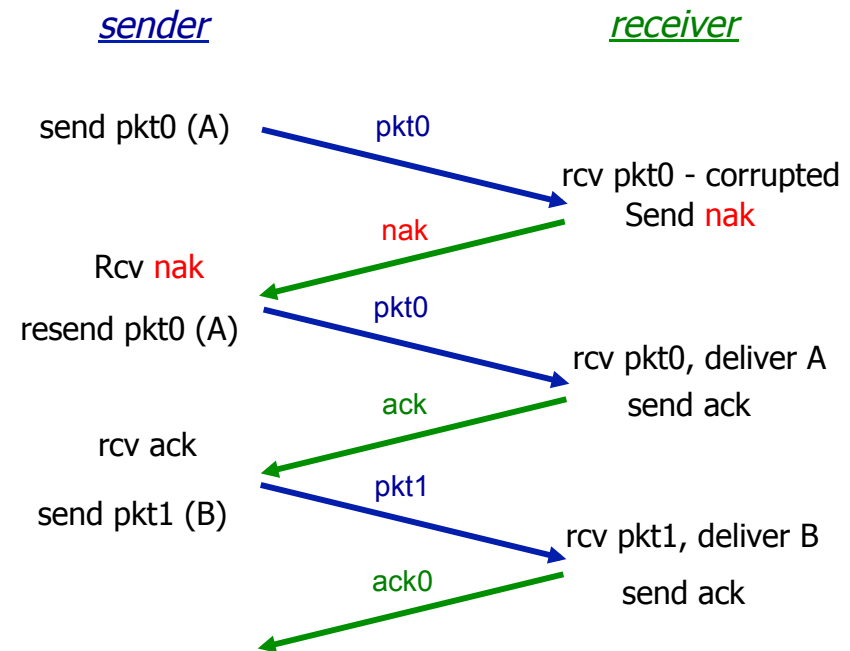


(a) No message or ACK corrupted

# Rdt2.1 in action - continued

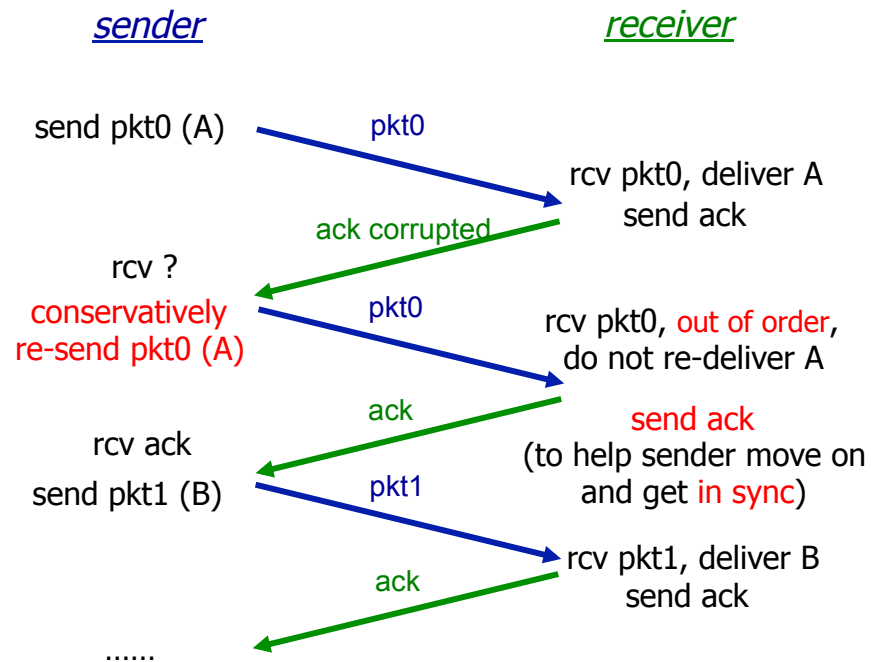


(a) No message or ACK corrupted

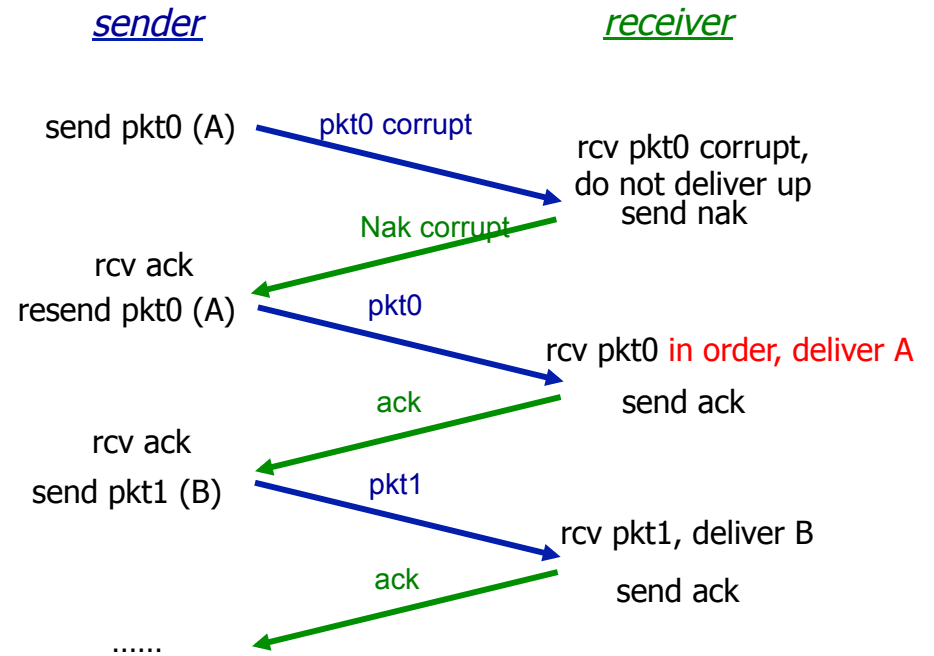


(b) Message corrupted, feedback ok

# Rdt2.1 in action - continued



(c) Message ok, ACK corrupted

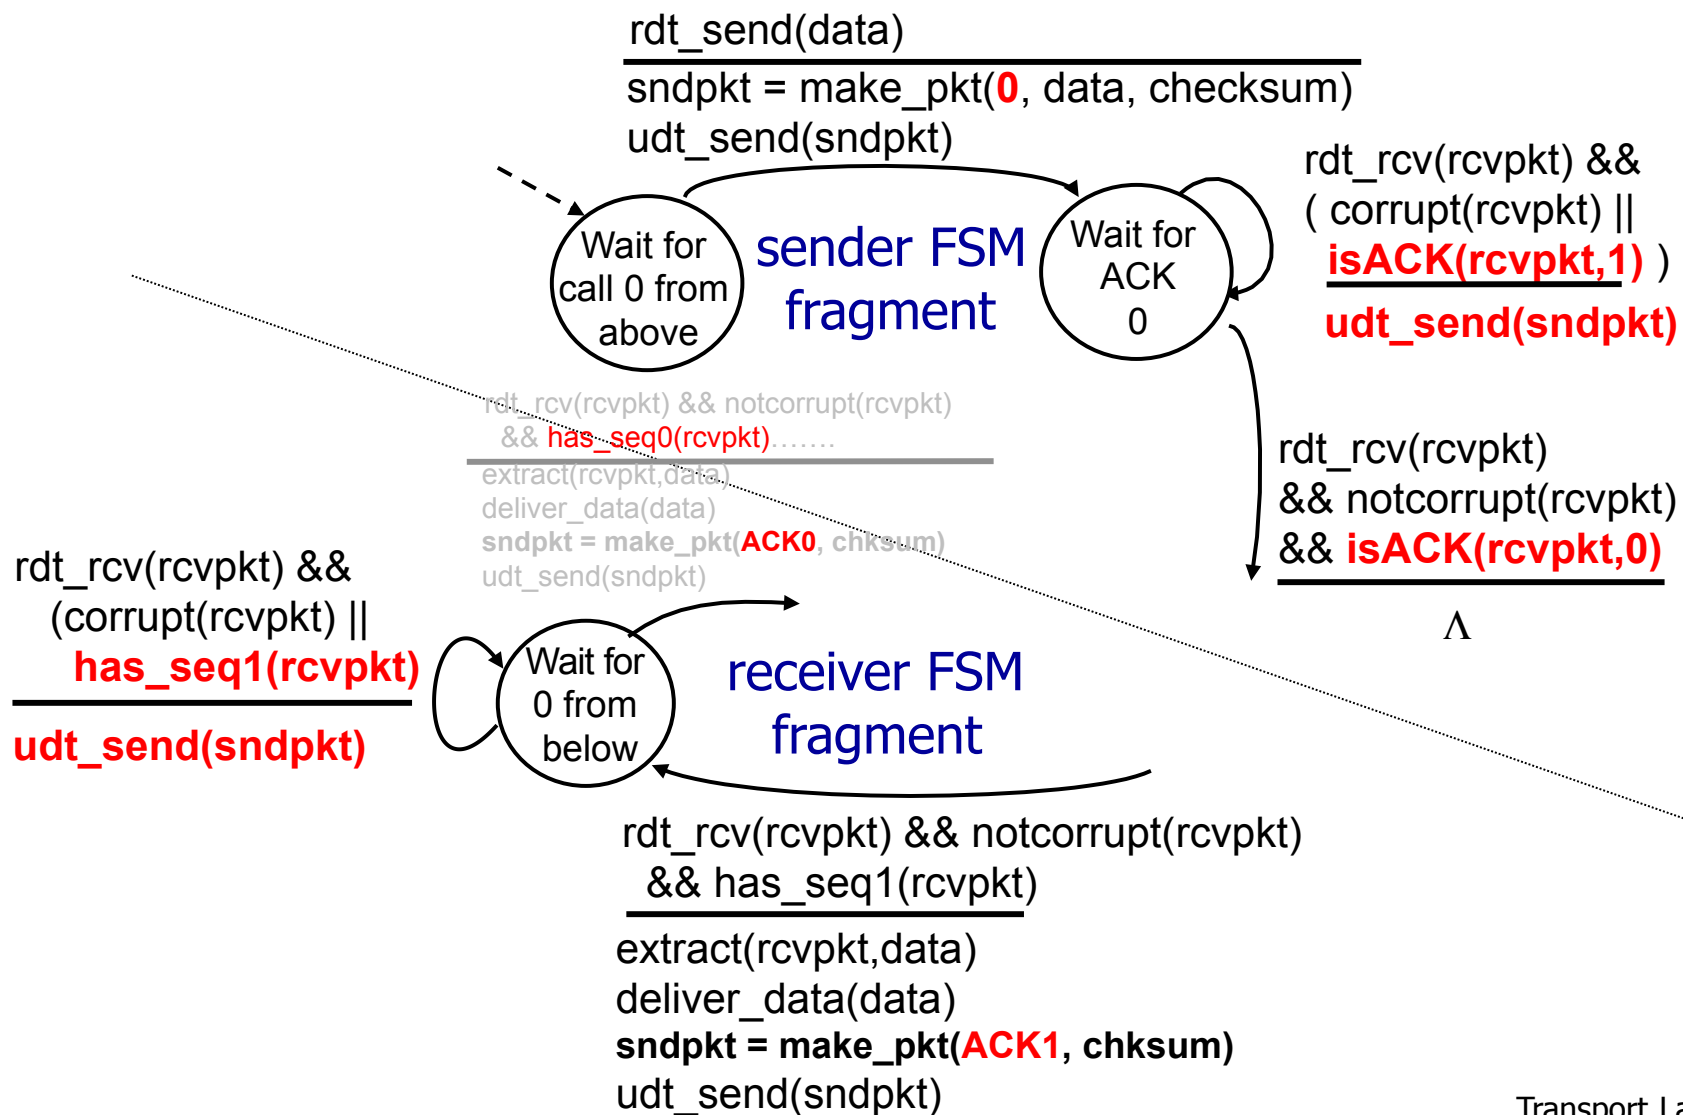


(d) Message corrupted, NAK corrupted

## rdt2.2: a NAK-free protocol

- ❖ same functionality as rdt2.1, using ACKs only
- ❖ instead of NAK, receiver sends ACK **for last pkt received OK** (=“duplicate ACK”)
  - receiver must **explicitly include seq # of pkt being ACKed**
- ❖ **duplicate ACK** at sender results in same action as NAK: *retransmit current pkt*

# rdt2.2: sender, receiver fragments



# rdt3.0: channels with errors *and* loss

## new assumption:

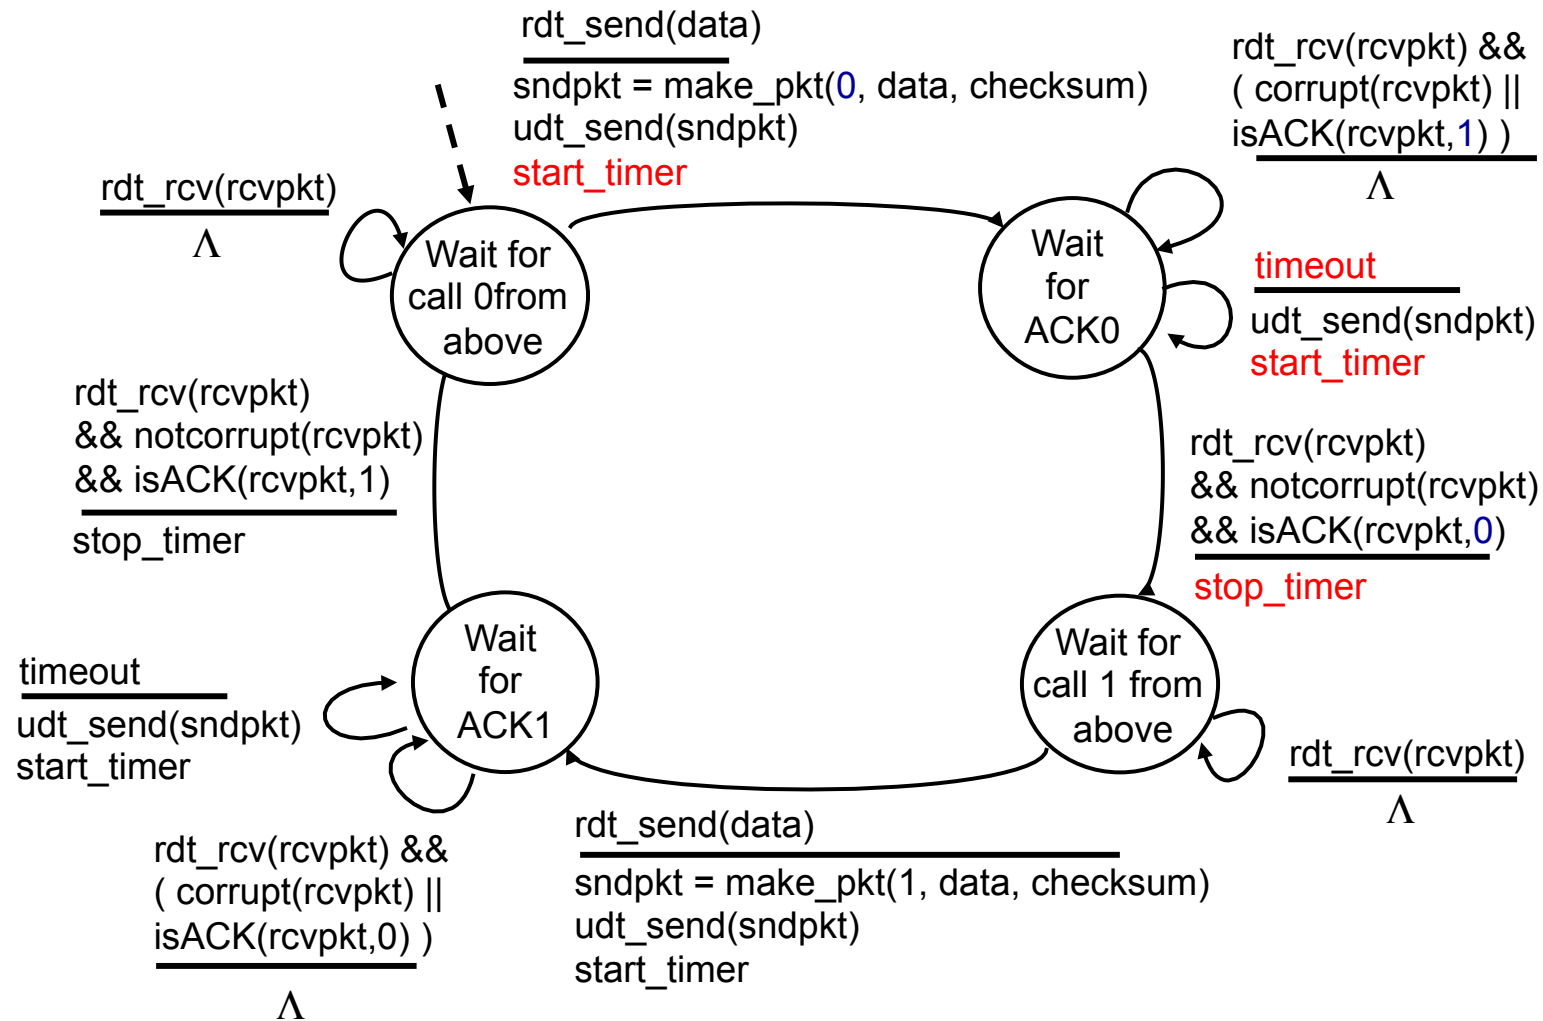
underlying channel can also lose packets (data, ACKs)

- How to detect loss?
- What to do in case of loss?
- checksum, seq. #, ACKs, retransmissions will be of help here, but not enough [why?]

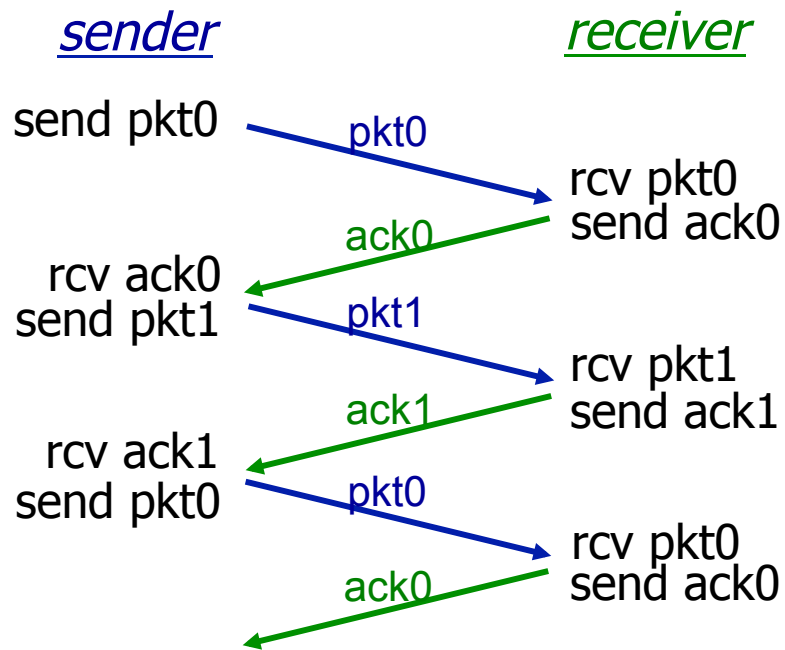
## approach: detect loss at server with Timeout

- ❖ Sender waits “reasonable” amount of time for ACK
- ❖ retransmits if no ACK received in this time
- ❖ if pkt (or ACK) just delayed (not lost):
  - retransmission will be duplicate, but seq. #'s already handles this
  - receiver must specify seq # of pkt being ACKed
- ❖ requires countdown timer


# rdt3.0 sender



# rdt3.0 in action



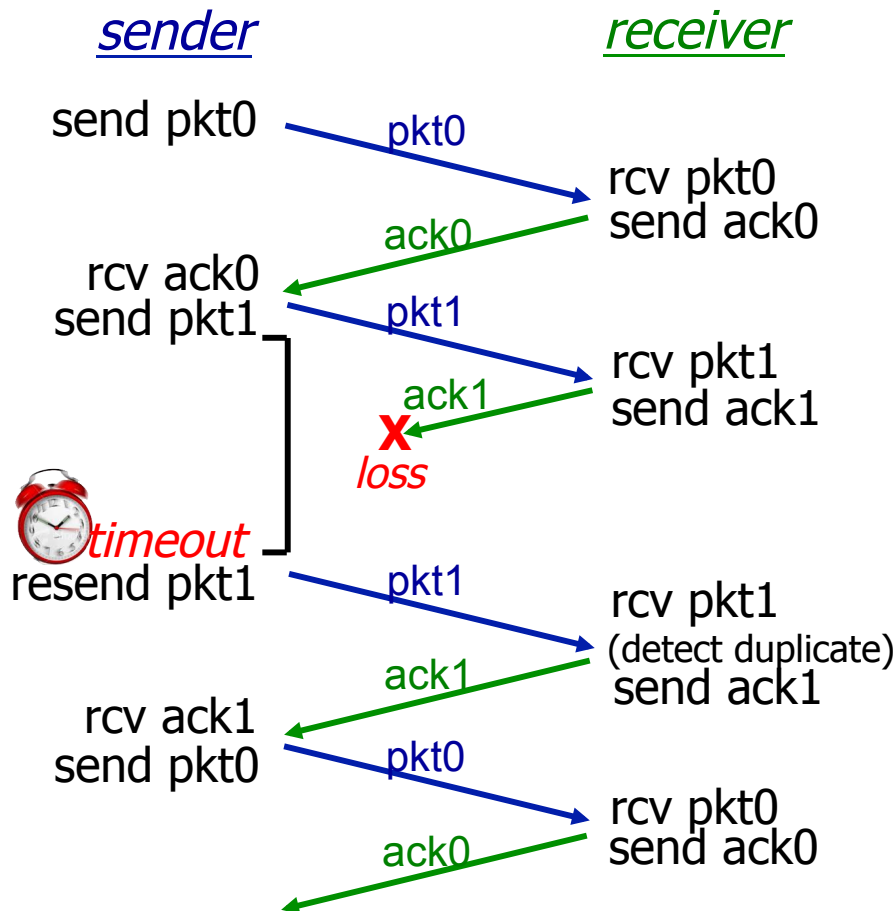
(a) no loss




(b) packet loss



# rdt3.0 in action



(c) ACK loss



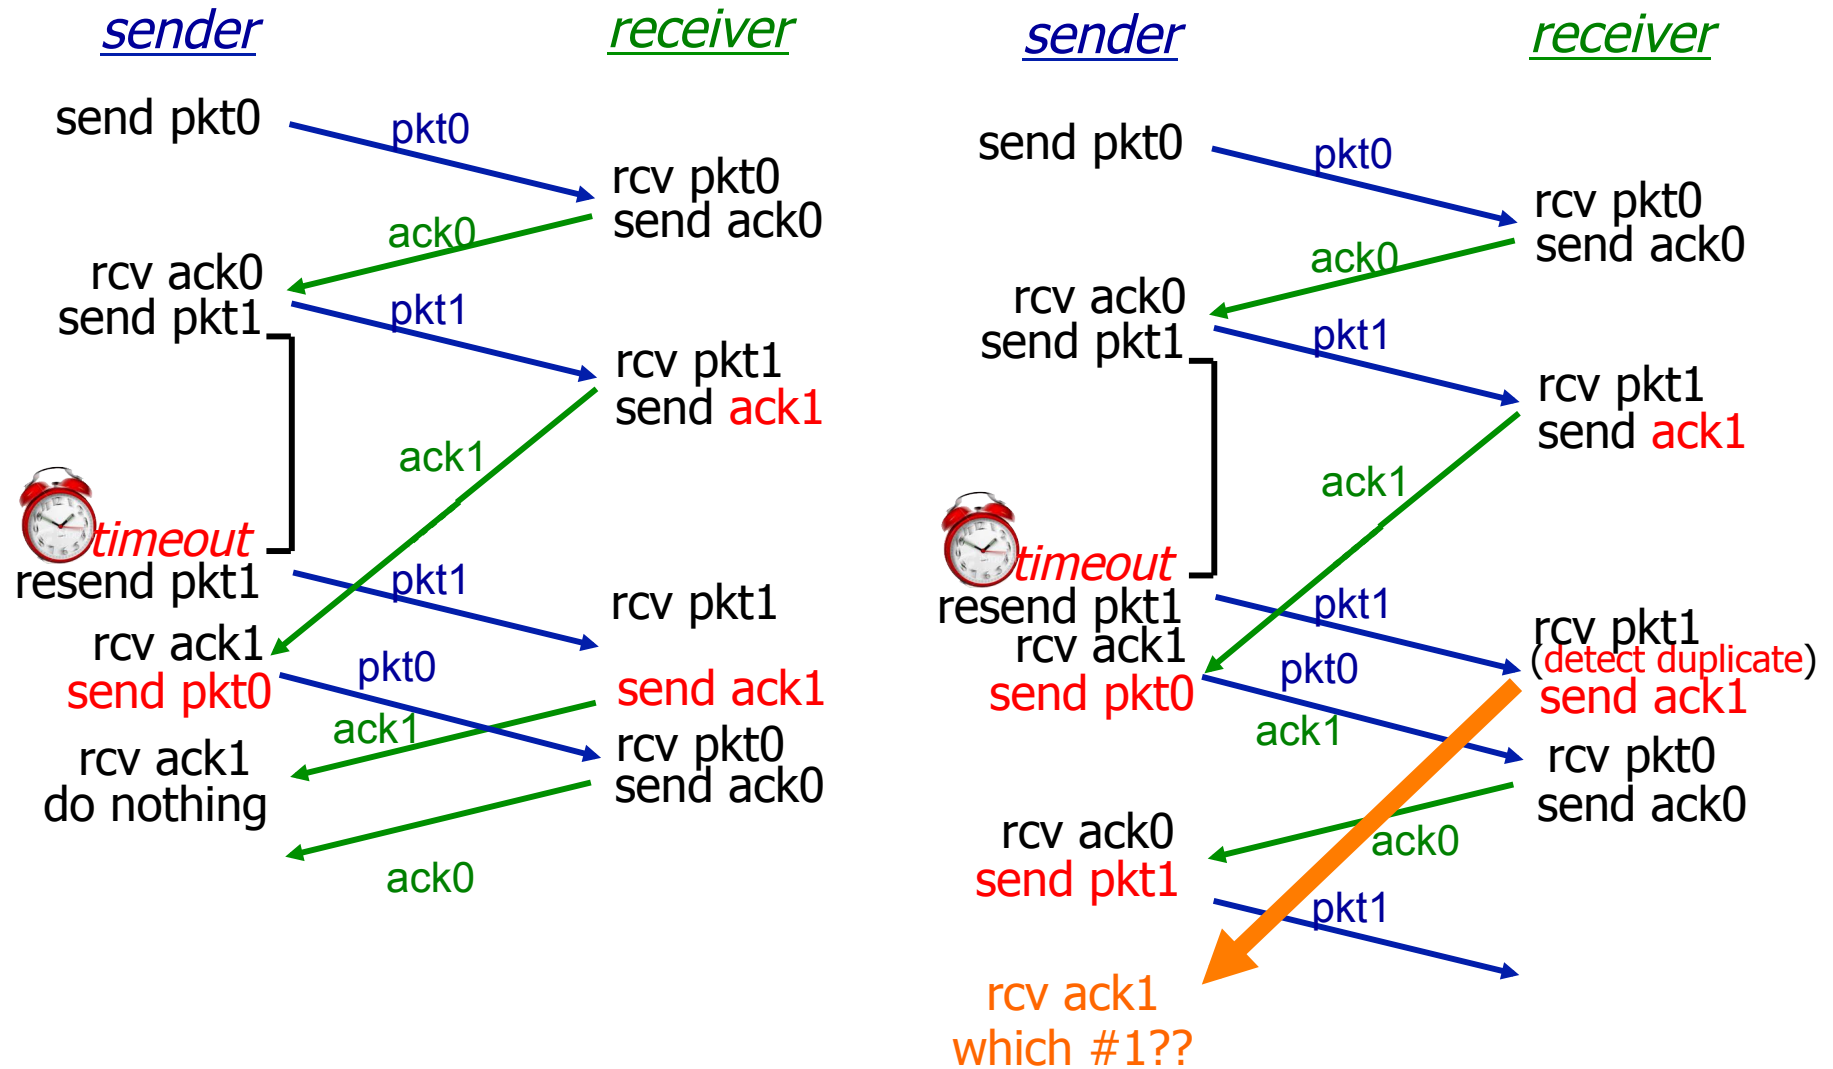
(d) premature timeout/ delayed ACK

# Summary of RDT mechanisms

---


- ❖ Corrupted bits (RDT 2.\*)
  - Detection via checksum
- ❖ Packet loss (RDT 3.\*)
  - Detection mechanisms
    - Feedback: ACK, NAK (NAK-free possible using duplicate ACKs)
    - No feedback: Timeout
  - Recovery mechanisms
    - Retransmission
    - Sequence Numbers (to detect duplicates): choose seq.no
  - Retransmission protocols
    - Window=1: stop-and-wait
    - Window>1: GBN, SR
    - Sender and receiver state can be out of sync.
- ❖ Reordering in the network?
  - Sequence numbers
  - Reordering looks like loss: may trigger retransmissions in TCP

# rdt3.0: delayed ack Q revisited



A: possible only if network reorders packets

# rdt3.0: stop-and-wait operation



$$U_{\text{sender}} = \frac{L / R}{RTT + L / R} = \frac{.008}{30.008} = 0.00027$$

# Performance of rdt3.0

- ❖ rdt3.0 is correct, but performance stinks
- ❖ e.g.: 1 Gbps link, 15 ms prop. delay, 8000 bit packet:

$$D_{trans} = \frac{L}{R} = \frac{8000 \text{ bits}}{10^9 \text{ bits/sec}} = 8 \text{ microseconds}$$

- $U_{\text{sender}}$  **utilization** (or “efficiency” in your HW): fraction of time sender busy sending


$$U_{\text{sender}} = \frac{L / R}{RTT + L / R} = \frac{.008}{30.008} = 0.00027$$

- if RTT=30 msec, 1KB pkt every 30 msec: 33kB/sec throughput over 1 Gbps link
  - Example of a (bad) network protocol limits use of physical resources (fast link)!

# Pipelined protocols

**pipelining:** sender allows multiple, “in-flight”, yet-to-be-acknowledged pkts

- range of sequence numbers must be increased
- buffering at sender and/or receiver




(a) a stop-and-wait protocol in operation

(b) a pipelined protocol in operation

- ❖ two generic forms of pipelined protocols: *go-Back-N*, *selective repeat*

# Pipelining: increased utilization



**Q:** What is the best choice of the “window” of pipelined messages?

**A:**  $RTT / (L/R + 1)$  keeps the channel always busy

# Pipelined protocols: overview

## Go-back-N:

- ❖ sender can have up to N unacked packets in pipeline
- ❖ receiver only sends *cumulative ack*
  - doesn't ack packet if there is a gap
- ❖ sender has **timer for oldest unacked packet**
  - when timer expires, retransmit *all* unacked packets


## Selective Repeat:

- ❖ sender can have up to N unack'ed packets in pipeline
- ❖ rcvr sends *individual ack* for each packet
- ❖ sender maintains **timer for each unacked packet**
  - when timer expires, retransmit only that unacked packet



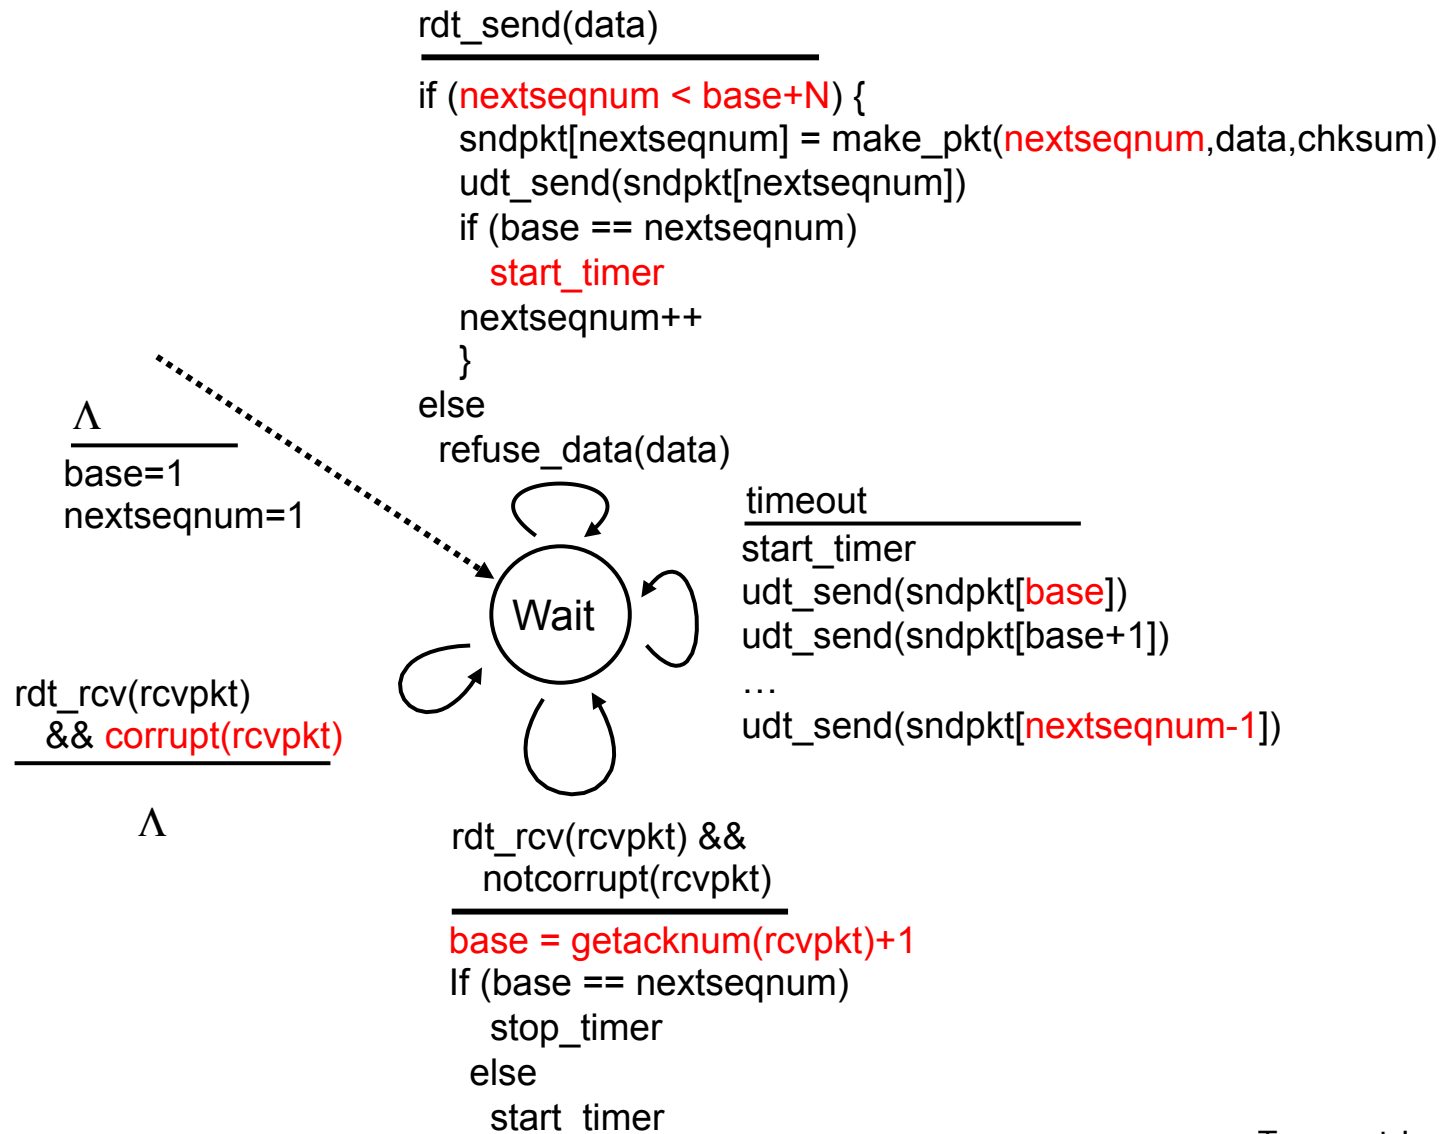
# Go-Back-N: sender

- ❖ k-bit sequence # in pkt header
- ❖ “sliding window” of up to N, consecutive unack’ed pkts allowed

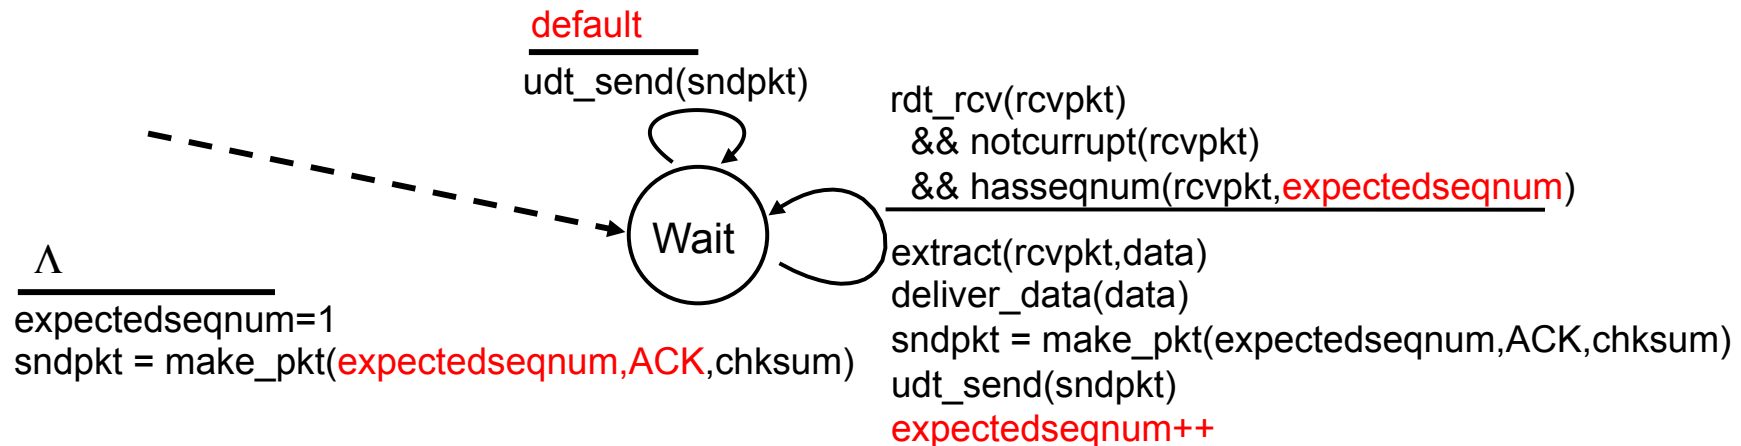


- ❖ Send\_base: oldest in-flight packet (n): first yellow packet
- ❖ ACK(m): ACKs all pkts up to, including seq # m - “cumulative ACK”
  - may receive duplicate ACKs (see receiver)
- ❖ timeout(n): keep timer for n and retransmit packet n and all higher seq # pkts in window – i.e., all yellow packets

# GBN: sender extended FSM




# GBN: receiver extended FSM



- ❖ Keeps track of correctly-received pkt with highest *in-order* seq #:
  - if sender-receiver in sync, this should be the last green packet in Sender's window
- ❖ ACK-only: always send ACK for that particular packet
  - may generate duplicate ACKs
  - need only remember one number **expectedseqnum**
    - If sender-receiver in sync, this should be the 1st yellow packet in Sender's window
- ❖ out-of-order pkt:
  - discard (don't buffer): *no receiver buffering! Why?*
  - re-ACK pkt with highest in-order seq #
- ❖ Really simple receiver!

# GBN in action



sender window (N=4)

0 1 2 3 4 5 6 7 8  
 0 1 2 3 4 5 6 7 8  
 0 1 2 3 4 5 6 7 8  
 0 1 2 3 4 5 6 7 8

0 1 2 3 4 5 6 7 8  
 0 1 2 3 4 5 6 7 8

0 1 2 3 4 5 6 7 8  
 0 1 2 3 4 5 6 7 8  
 0 1 2 3 4 5 6 7 8  
 0 1 2 3 4 5 6 7 8

0 1 2 3 4 5 6 7 8

sender

send pkt0  
 send pkt1  
 send pkt2  
 send pkt3  
 (wait)

rcv ack0, send pkt4  
 rcv ack1, send pkt5

ignore duplicate ACK



*pkt 2 timeout*

send pkt2  
 send pkt3  
 send pkt4  
 send pkt5

rcv ack2

receiver

receive pkt0, send ack0  
 receive pkt1, send ack1

receive pkt3, **discard**,  
 (re)send ack1

receive pkt4, **discard**,  
 (re)send ack1

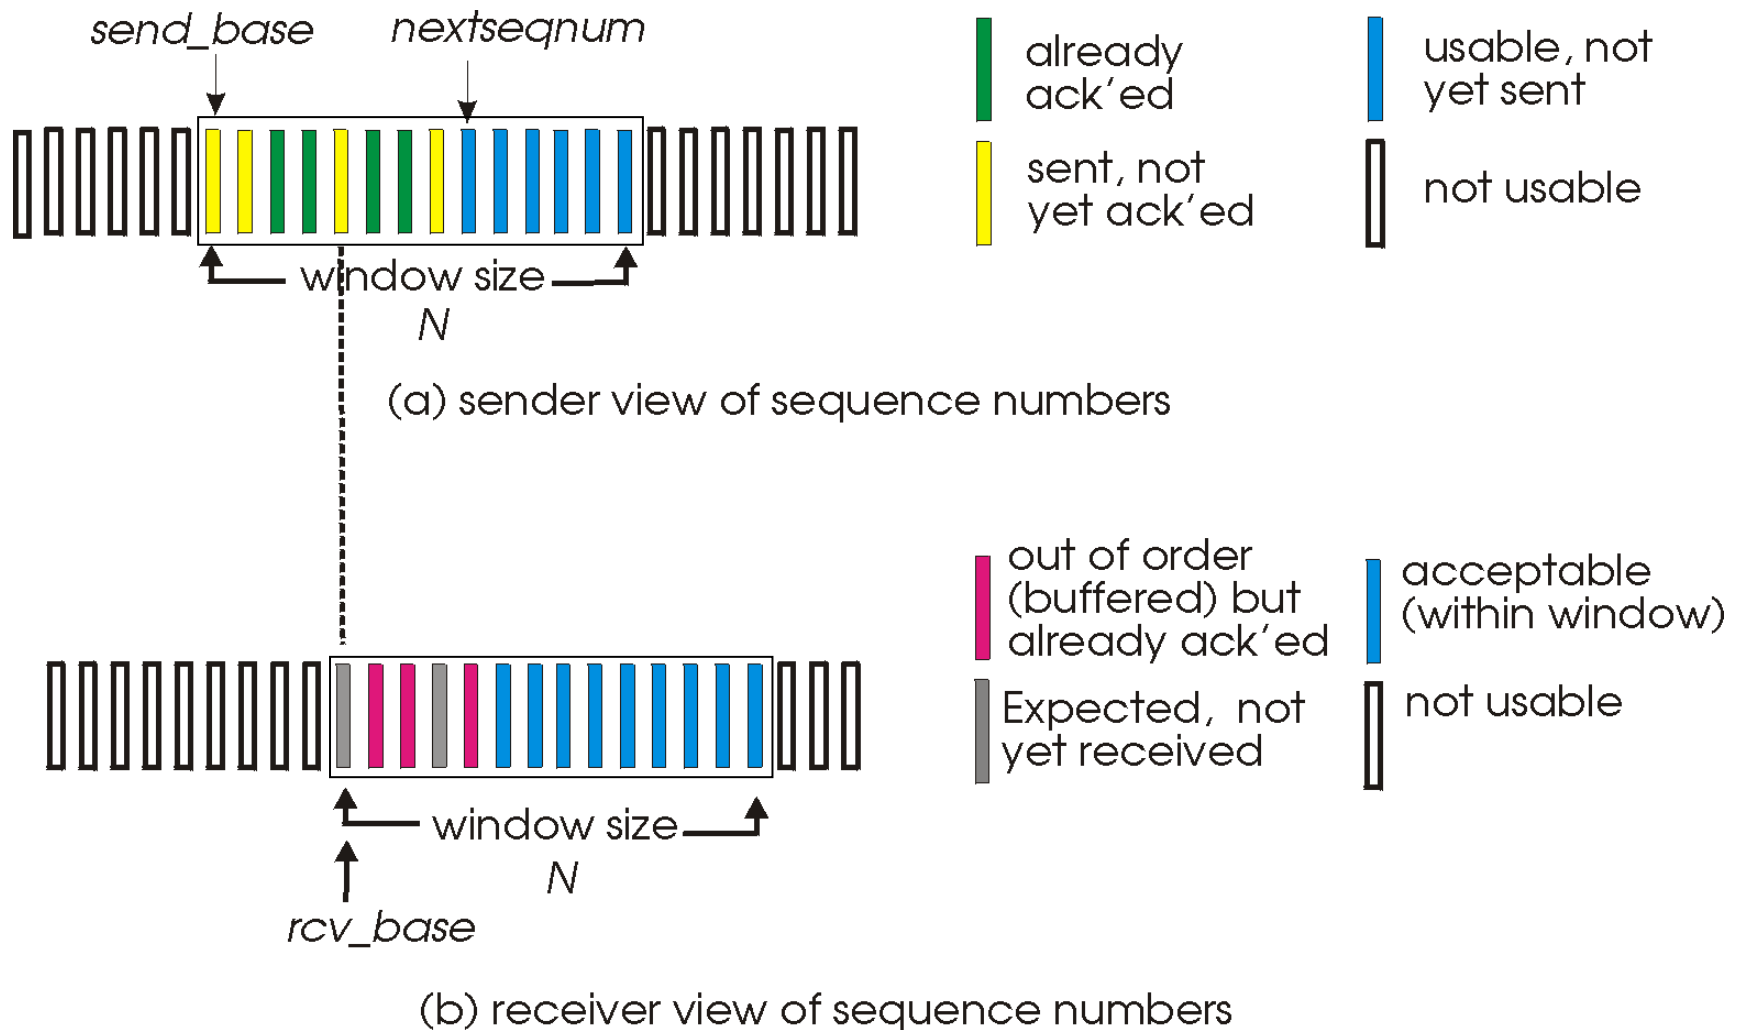
receive pkt5, **discard**,  
 (re)send ack1

rcv pkt2, deliver, **send ack2**  
 rcv pkt3, deliver, send ack3  
 rcv pkt4, deliver, send ack4  
 rcv pkt5, deliver, send ack5

# Selective Repeat (SR)

- ❖ Receiver *individually acknowledges all* correctly received pkts
  - whether they are received in-order or out-of-order
  - buffers pkts, as needed, for eventual in-order delivery to upper layer
- ❖ Sender *only retransmits un-ACKed pkts*
  - sender maintains timer for each unACKed pkt
- ❖ Both *sender and receiver* maintain *windows*
  - Sender window
    - $N$  consecutive seq #'s
    - limits seq #'s of sent, unACKed pkts
  - Receiver window
    - $N$  consecutive seq #'s
    - Limits seq#'s of buffered out-of-order packets

# Selective repeat: sender, receiver windows



# Selective Repeat (FSM not shown)

## sender

### data from above:

- ❖ if next available seq # in window, send pkt

### timeout(n):

- ❖ resend pkt n, restart timer

### ACK(n) in [sendbase, sendbase+N]:

- ❖ mark pkt n as received
- ❖ if n smallest unACKed pkt, advance window base to next unACKed seq #

## receiver

### pkt n in [rcvbase, rcvbase+N-1]


- ❖ send ACK(n)
- ❖ out-of-order: buffer
- ❖ in-order: deliver (also deliver all buffered, in-order pkts), advance window to next not-yet-received pkt

### pkt n in [rcvbase-N, rcvbase-1]


- ❖ ACK(n) [Why not ignore?]

### otherwise:

- ❖ ignore



(a) sender view of sequence numbers



(b) receiver view of sequence numbers

# Selective repeat in action

sender window (N=4)

0 1 2 3 4 5 6 7 8  
 0 1 2 3 4 5 6 7 8  
 0 1 2 3 4 5 6 7 8  
 0 1 2 3 4 5 6 7 8  
 [empty]

0 1 2 3 4 5 6 7 8  
 0 1 2 3 4 5 6 7 8

0 1 2 3 4 5 6 7 8  
 0 1 2 3 4 5 6 7 8  
 0 1 2 3 4 5 6 7 8  
 0 1 2 3 4 5 6 7 8

sender

send pkt0  
 send pkt1  
 send pkt2  
 send pkt3  
 (wait)

rcv ack0, send pkt4  
 rcv ack1, send pkt5

record ack3 arrived



*pkt 2 timeout*

send pkt2

record ack4 arrived

record ack4 arrived

*Q: what happens when ack2 arrives?*

receiver

receive pkt0, send ack0  
 receive pkt1, send ack1

receive pkt3, buffer,  
 send ack3

receive pkt4, buffer,  
 send ack4

receive pkt5, buffer,  
 send ack5

rcv pkt2; deliver pkt2,  
 pkt3, pkt4, pkt5; send ack2




## Lack of sync + finite seq#: ambiguity

### Example:


- ❖ seq #'s: 0, 1, 2, 3
- ❖ window size=3
- ❖ receiver sees no difference in two scenarios!
- ❖ duplicate data accepted as new in (b)

**Q:** what relationship between seq # size (SN) and window size (N) to avoid problem?

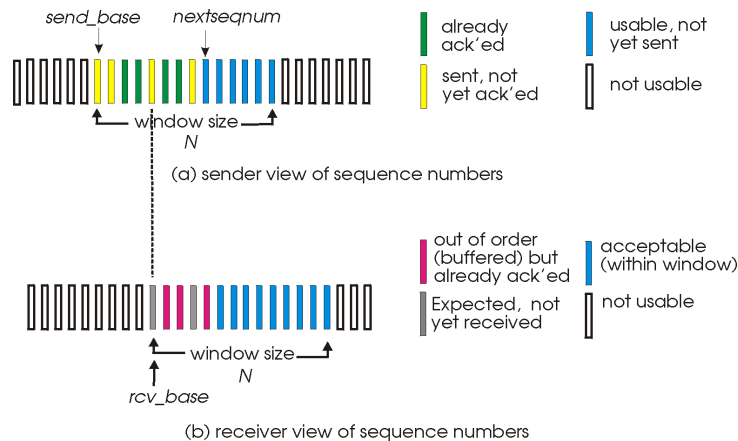
**A:**  $SN \geq 2N$  because sender and receiver window must have overlap of at least 1



receiver can't see sender side.  
receiver behavior identical in both cases!  
*something's (very) wrong!*




# SR: sender, receiver windows & Seq No



N= Sender window  
= Receiver window

←-----→ SN: sequence numbers for packets and acks



Sender & receiver windows overlap. At most consecutive.  
 $SN \geq 2N$ , so that seqnums for packets and acks are unique



Sender & receiver windows in perfect sync



Impossible



Impossible



Special Case: Stop and Wait: Sender  $W=2$ , Receiver  $W=1$

# Go-back-N vs SR: Mechanisms

## Go-back-N:

- ❖ Sender can have up to N unack'ed packets in pipeline
  - sliding window
- ❖ Rcvr sends *cumulative* ack for last in-order packet
  - maintains *expectedseqnum*
  - doesn't accept or ack out-of-order packet
- ❖ Sender maintains timer for *oldest unacked* packet
  - if timer expires, *retransmit all* unack'ed packets

## Selective Repeat:

- ❖ Sender can have up to N unack'ed packets in pipeline
  - sliding window
- ❖ Rcvr sends *individual ack* for each packet
  - maintains *Rcvr window*
  - buffers and acks all packets within Rcvr window
- ❖ Sender maintains timer for *each unacked* packet
  - when timer expires, *retransmit only that one* unack'ed packet

## GBN vs SR: Performance

- ❖ Compared to Stop-and Wait
  - They both fill the pipeline
- ❖ Loss rate
  - Light loss:
    - SR: selectively retransmits what is needed
    - GBN: a single packet lost causes unnecessary retransmission of all packets in the window
  - Heavy loss
    - GBN: ok
- ❖ Complexity:
  - GBN is simpler – less state

# Practice GBN vs SR

## ❖ Sample Midterm

### ■ Problem 4

- <https://eee.uci.edu/16s/18105/hws/eecs148-midterm-s15.pdf>

## ❖ Companion Website

### ■ [http://wps.pearsoned.com/ecs\\_kurose\\_compnetw\\_6/](http://wps.pearsoned.com/ecs_kurose_compnetw_6/)

### ■ Applets for Ch.3

- [http://wps.pearsoned.com/ecs\\_kurose\\_compnetw\\_6/216/55463/14198700.cw/index.html](http://wps.pearsoned.com/ecs_kurose_compnetw_6/216/55463/14198700.cw/index.html)

### ■ interactive exercises for Ch.3 #2

- [http://wps.pearsoned.com/ecs\\_kurose\\_compnetw\\_6/216/55463/14198700.cw/index.html](http://wps.pearsoned.com/ecs_kurose_compnetw_6/216/55463/14198700.cw/index.html)